Synthetic example. *Fictional organization and self-reported answers generated to demonstrate AEGIS outputs. Not evidence about any real organization.*

AI Governance Maturity Report — Example Bank plc (synthetic)

Assessment: AI Governance Maturity Self-Assessment (template AEGIS-SA-MATURITY v1.0.0)
Generated: 2026-10-21
Evidence status: User-asserted — all answers are self-reported; no evidence has been reviewed.

This self-assessment produces an indicative, self-reported view of AI governance maturity. All results are User-asserted: no evidence has been reviewed. Framework readiness indicators are not a determination of conformity with ISO/IEC 42001, compliance with the EU AI Act, or alignment with NIST AI RMF. AEGIS does not certify compliance and this report is not legal advice.

1. Overall maturity

Overall score	Maturity level	Interpretation
56.4 / 100	**Developing**	Basic practices exist but are inconsistent and hard to evidence.

2. Domain scores

Domain	Weight	Score	Profile
Strategy & Accountability	12%	66.7	[##############.......]
AI Inventory & Visibility	14%	46.7	[#########..........]
Risk Management & Lifecycle	14%	60.0	[############........]
Data & Privacy	12%	60.0	[############........]
Security & Resilience	14%	46.7	[#########..........]
Fairness, Safety & Transparency	12%	66.7	[##############.......]
Human Oversight & Operations	12%	53.3	[###########.........]
Assurance, Evidence & Improvement	10%	53.3	[###########.........]

Scoring: domain score = sum of answers ÷ maximum possible × 100; overall = Σ(domain score × weight). Source: AEGIS self-assessment template AEGIS-SA-MATURITY v1.0.0.

3. Top strengths

- **Risk Management & Lifecycle** — Do AI use cases undergo risk assessment before production deployment? (answer 4/5; question RSK-Q1; controls GL-01, GL-02, GL-03) — *User-asserted*
- **Strategy & Accountability** — Does your organization have a documented AI governance policy and accountable ownership model? (answer 4/5; question STR-Q1; controls AO-01, AO-03) — *User-asserted*
- **Data & Privacy** — Are privacy, consent, retention, cross-border transfer, and data-minimization requirements assessed? (answer 4/5; question DAT-Q2; controls IS-03) — *User-asserted*
- **Fairness, Safety & Transparency** — Are AI systems assessed for bias, discriminatory outcomes, safety risks, and unintended harm? (answer 4/5; question FST-Q1; controls IS-05, IS-06) — *User-asserted*

4. Top gaps

#	Domain	Question	Answer	Recommended action	AEGIS controls
1	AI Inventory & Visibility	INV-Q1: Does your organization maintain a complete inventory of AI systems, models, agents, and vendor AI products?	2/5	Stand up an AI system of record and run inventory workshops plus shadow-AI discovery to reach full coverage.	EI-01, EI-08, EI-09

#	Domain	Question	Answer	Recommended action	AEGIS controls
2	AI Inventory & Visibility	INV-Q3: Can you identify the owner, purpose, users, data sources, and dependencies for each material AI system?	2/5	Complete owner, purpose, users, data-source and dependency fields for all material AI systems.	EI-03, EI-04, EI-06, EI-10
3	Risk Management & Lifecycle	RSK-Q3: Is there a defined process for retiring, restricting, or remediating AI systems that no longer meet requirements?	2/5	Define restrict/remediate/retire procedures and an exception register with expiry dates.	GL-09, GL-10
4	Security & Resilience	SEC-Q2: Are LLM applications and agents tested for prompt injection, jailbreak, data leakage, and unsafe tool use?	2/5	Run adversarial testing (prompt injection, jailbreak, leakage, unsafe tool use) before release and after material change.	IS-14, IS-11, IS-12
5	Security & Resilience	SEC-Q3: Can you disable or roll back an AI system, model, agent, or tool quickly after a security or safety incident?	2/5	Implement and test a kill switch and rollback runbook for every production AI system and agent.	IS-15

5. Critical risk flags

- **Critical gap — No reliable AI inventory** (trigger: INV-Q1). Without a complete inventory, risk tiering, oversight and evidence cannot be applied to systems you do not know about. Controls: EI-01, EI-09. Evidence status: User-asserted.
- **Critical gap — Sensitive data may reach prompts, logs or external models** (trigger: DAT-Q3). Unredacted personal or confidential data in prompts and logs creates privacy and confidentiality exposure. Controls: IS-04, IS-03. Evidence status: User-asserted.
- **Critical gap — Agents or LLM apps not tested for prompt injection and unsafe tool use** (trigger: SEC-Q2). Agents with tool access that are untested for injection can be manipulated into unauthorised actions or data leakage. Controls: IS-14, IS-11, IS-12. Evidence status: User-asserted.
- **Critical gap — No tested kill switch or rollback** (trigger: SEC-Q3). Without containment, an AI incident cannot be stopped quickly. Controls: IS-15. Evidence status: User-asserted.
- **Critical gap — Cannot produce audit-ready evidence** (trigger: ASR-Q2). Inability to evidence controls undermines responses to auditors, customers and regulators. Controls: SE-05, SE-08. Evidence status: User-asserted.

6. 30/60/90-day action plan

First 30 days — stabilise

- Stand up an AI system of record and run inventory workshops plus shadow-AI discovery to reach full coverage. (AI Inventory & Visibility; INV-Q1; controls EI-01, EI-08, EI-09)
- Complete owner, purpose, users, data-source and dependency fields for all material AI systems. (AI Inventory & Visibility; INV-Q3; controls EI-03, EI-04, EI-06, EI-10)
- Define restrict/remediate/retire procedures and an exception register with expiry dates. (Risk Management & Lifecycle; RSK-Q3; controls GL-09, GL-10)
- Run adversarial testing (prompt injection, jailbreak, leakage, unsafe tool use) before release and after material change. (Security & Resilience; SEC-Q2; controls IS-14, IS-11, IS-12)
- Implement and test a kill switch and rollback runbook for every production AI system and agent. (Security & Resilience; SEC-Q3; controls IS-15)

Days 31–60 — standardise

- Apply AEGIS risk-tier criteria to every inventory record and approve tier assignments. (AI Inventory & Visibility; INV-Q2; controls EI-02)

- Add AI-specific change triggers (model, prompt, data, vendor, tool, agent) to change management with reassessment rules. (Risk Management & Lifecycle; RSK-Q2; controls GL-04, EI-05)
- Extend threat modelling, security review and vulnerability management to all AI systems and model supply chains. (Security & Resilience; SEC-Q1; controls IS-01, IS-10)
- Charter a cross-functional AI governance forum with documented authority to approve, restrict or stop use cases. (Strategy & Accountability; STR-Q2; controls AO-02)
- Publish an AI RACI and escalation matrix covering business, technology, risk, legal, privacy and security. (Strategy & Accountability; STR-Q3; controls AO-03, AO-06)

Days 61–90 — scale and automate

- Make a tier-proportionate risk and impact assessment a mandatory gate before any production deployment. (Risk Management & Lifecycle; RSK-Q1; controls GL-01, GL-02, GL-03)
- Approve and publish an AI governance policy with a named executive owner and accountability model. (Strategy & Accountability; STR-Q1; controls AO-01, AO-03)
- Embed privacy impact assessment (consent, retention, transfers, minimisation) into AI intake and change. (Data & Privacy; DAT-Q2; controls IS-03)
- Introduce bias, safety and harm evaluation for high-impact systems with documented thresholds. (Fairness, Safety & Transparency; FST-Q1; controls IS-05, IS-06)

7. Recommended AEGIS controls

Control	Pillar	Priority	NIST AI RMF	ISO/IEC 42001	EU AI Act
EI-01 Enterprise AI inventory (system of record)	EI	critical	GOVERN 1.6	A.4.2, 7.5	Art. 26, Art. 49
EI-09 Shadow AI discovery	EI	medium	GOVERN 1.6	A.4.2	—
IS-04 Sensitive data redaction and masking	IS	critical	MEASURE 2.10	A.7.4, A.6.2.6	Art. 10
IS-03 Privacy assessment and data minimisation	IS	critical	MEASURE 2.10, MAP 1.1	A.7.4, A.5.4	Art. 10
IS-14 Prompt-injection and data-leakage testing	IS	critical	MEASURE 2.7	A.6.2.4	Art. 15
IS-11 Least-privilege tool access	IS	critical	MEASURE 2.7, MANAGE 2.2	A.6.2.6, A.4.4	Art. 15, Art. 14
IS-12 Runtime guardrails and policy enforcement	IS	high	MANAGE 2.2, MANAGE 2.4, MEASURE 2.7	A.6.2.6	Art. 15, Art. 14
IS-15 Kill switch and rollback	IS	critical	MANAGE 2.4	A.6.2.6	Art. 14
SE-05 Audit-ready evidence packs	SE	high	GOVERN 1.4, MEASURE 2.9	7.5, 9.2	Art. 11, Art. 18
SE-08 Immutable action logging	SE	critical	MEASURE 2.8, MANAGE 4.1	A.6.2.8	Art. 12, Art. 19, Art. 26

Framework mappings are indicative AEGIS interpretations (Control Library v1.0.0).

8. Framework readiness indicators (indicative — not a compliance determination)

NIST AI RMF

Area	Indicator	Band	Based on questions
GOVERN	60.0	Partial readiness	STR-Q1, STR-Q2, STR-Q3, HUM-Q3, ASR-Q3
MAP	56.0	Partial readiness	INV-Q1, INV-Q2, INV-Q3, RSK-Q1, DAT-Q1
MEASURE	64.0	Partial readiness	SEC-Q2, FST-Q1, HUM-Q2, ASR-Q1, DAT-Q2
MANAGE	48.0	Limited readiness	RSK-Q2, RSK-Q3, SEC-Q3, HUM-Q1, ASR-Q2

ISO/IEC 42001

Area	Indicator	Band	Based on questions
Context & leadership (Cl. 4–5, A.2–A.3)	66.7	Partial readiness	STR-Q1, STR-Q2, STR-Q3
Planning & impact assessment (Cl. 6, A.5)	73.3	Partial readiness	INV-Q2, RSK-Q1, FST-Q1
Support & documented information (Cl. 7, A.4)	40.0	Limited readiness	HUM-Q3, INV-Q1, ASR-Q2
Operation & lifecycle (Cl. 8, A.6–A.7)	65.0	Partial readiness	RSK-Q2, DAT-Q1, DAT-Q2, SEC-Q1
Performance evaluation (Cl. 9)	60.0	Partial readiness	HUM-Q2, ASR-Q1
Improvement (Cl. 10)	50.0	Partial readiness	ASR-Q3, RSK-Q3

EU AI Act

Area	Indicator	Band	Based on questions
AI literacy (Art. 4)	40.0	Limited readiness	HUM-Q3
Risk management (Art. 9)	66.7	Partial readiness	RSK-Q1, RSK-Q2, INV-Q2
Data governance (Art. 10)	73.3	Partial readiness	DAT-Q1, DAT-Q2, FST-Q1
Technical documentation & record-keeping (Art. 11–12)	40.0	Limited readiness	INV-Q3, ASR-Q2
Transparency (Art. 13, Art. 50)	60.0	Partial readiness	FST-Q2, FST-Q3
Human oversight (Art. 14)	50.0	Partial readiness	HUM-Q1, SEC-Q3
Accuracy, robustness & cybersecurity (Art. 15)	50.0	Partial readiness	SEC-Q1, SEC-Q2
Post-market monitoring & incidents (Art. 72–73)	60.0	Partial readiness	HUM-Q2, ASR-Q1

These indicators summarise self-reported practices that relate to each framework area. They do not indicate conformity with ISO/IEC 42001, compliance with the EU AI Act, or alignment with NIST AI RMF, and EU AI Act applicability depends on your role (provider/deployer) and system classification.

9. Recommended AEGIS modules

- **AEGIS Assess** — Establish an evidence-based baseline of maturity and use-case risk.
- **AEGIS Register** — Create a single, owned and classified AI system of record.
- **AEGIS Shield** — Enforce runtime guardrails, approvals and containment for LLM apps and agents.
- **AEGIS Assurance** — Test controls continuously and produce audit-ready evidence packs.
- **AEGIS Agent** — Accelerate assessments, gap analysis and reporting with a citation-backed governance copilot.

10. Next step: AEGIS Governance Baseline

A fixed-scope 4–6 week engagement — inventory up to 25 AI systems, validate this maturity baseline with evidence, deep-dive five priority systems, review two agents, and deliver a gap register, remediation roadmap and executive readiness report.

Book an AEGIS Governance Baseline: /contact/book-governance-baseline/

Appendix A — Your answers

Question	Answer
STR-Q1: Does your organization have a documented AI governance policy and accountable ownership model?	4 — Consistently implemented
STR-Q2: Is there a cross-functional AI governance forum with authority to approve, restrict, or stop AI use cases?	3 — Defined for some systems
STR-Q3: Are AI roles, responsibilities, and escalation paths defined for business, technology, risk, legal, and security teams?	3 — Defined for some systems
INV-Q1: Does your organization maintain a complete inventory of AI systems, models, agents, and vendor AI products?	2 — Initial / ad hoc

Question	Answer
INV-Q2: Are AI systems classified by business criticality, data sensitivity, decision impact, autonomy, and regulatory relevance?	3 — Defined for some systems
INV-Q3: Can you identify the owner, purpose, users, data sources, and dependencies for each material AI system?	2 — Initial / ad hoc
RSK-Q1: Do AI use cases undergo risk assessment before production deployment?	4 — Consistently implemented
RSK-Q2: Are model, prompt, data, vendor, and agent changes subject to change control and reassessment?	3 — Defined for some systems
RSK-Q3: Is there a defined process for retiring, restricting, or remediating AI systems that no longer meet requirements?	2 — Initial / ad hoc
DAT-Q1: Are data sources, data classifications, and data flows documented for AI systems?	3 — Defined for some systems
DAT-Q2: Are privacy, consent, retention, cross-border transfer, and data-minimization requirements assessed?	4 — Consistently implemented
DAT-Q3: Are sensitive data fields redacted, masked, or excluded from prompts, logs, and external model calls?	2 — Initial / ad hoc
SEC-Q1: Are AI systems included in threat modeling, security review, and vulnerability-management processes?	3 — Defined for some systems
SEC-Q2: Are LLM applications and agents tested for prompt injection, jailbreak, data leakage, and unsafe tool use?	2 — Initial / ad hoc
SEC-Q3: Can you disable or roll back an AI system, model, agent, or tool quickly after a security or safety incident?	2 — Initial / ad hoc
FST-Q1: Are AI systems assessed for bias, discriminatory outcomes, safety risks, and unintended harm?	4 — Consistently implemented
FST-Q2: Are users informed when they are interacting with AI or when AI materially influences a decision?	3 — Defined for some systems
FST-Q3: Are explanations, disclosures, and appeal mechanisms available for affected individuals where appropriate?	3 — Defined for some systems
HUM-Q1: Are human-review, override, escalation, and appeal mechanisms defined for high-impact AI decisions?	3 — Defined for some systems
HUM-Q2: Are AI systems monitored for performance, drift, misuse, incidents, and unexpected behavior?	3 — Defined for some systems
HUM-Q3: Are operators and affected teams trained on AI limitations, escalation paths, and responsible use?	2 — Initial / ad hoc
ASR-Q1: Are AI controls tested periodically, with evidence retained for audit and review?	3 — Defined for some systems
ASR-Q2: Can your organization produce an audit-ready evidence pack for a material AI system on request?	2 — Initial / ad hoc
ASR-Q3: Are AI governance metrics reviewed by leadership and used to drive corrective action?	3 — Defined for some systems

Appendix B — Sources

- AEGIS Control EI-01: Enterprise AI inventory (system of record) (Control Library v1.0.0)
- AEGIS Control EI-09: Shadow AI discovery (Control Library v1.0.0)
- AEGIS Control IS-04: Sensitive data redaction and masking (Control Library v1.0.0)
- AEGIS Control IS-03: Privacy assessment and data minimisation (Control Library v1.0.0)
- AEGIS Control IS-14: Prompt-injection and data-leakage testing (Control Library v1.0.0)
- AEGIS Control IS-11: Least-privilege tool access (Control Library v1.0.0)
- AEGIS Control IS-12: Runtime guardrails and policy enforcement (Control Library v1.0.0)
- AEGIS Control IS-15: Kill switch and rollback (Control Library v1.0.0)
- AEGIS Control SE-05: Audit-ready evidence packs (Control Library v1.0.0)
- AEGIS Control SE-08: Immutable action logging (Control Library v1.0.0)

- AEGIS Self-Assessment template AEGIS-SA-MATURITY v1.0.0

AEGIS supports governance readiness, control implementation and evidence collection. It does not certify ISO/IEC 42001 conformity, guarantee EU AI Act compliance, or replace legal advice. Framework mappings are indicative.