Synthetic example. *Fictional organization and self-reported answers generated to demonstrate AEGIS outputs. Not evidence about any real organization.*

AI Governance Maturity Report — Example Health Group (synthetic)

Assessment: AI Governance Maturity Self-Assessment (template AEGIS-SA-MATURITY v1.0.0)
Generated: 2026-10-21
Evidence status: User-asserted — all answers are self-reported; no evidence has been reviewed.

This self-assessment produces an indicative, self-reported view of AI governance maturity. All results are User-asserted: no evidence has been reviewed. Framework readiness indicators are not a determination of conformity with ISO/IEC 42001, compliance with the EU AI Act, or alignment with NIST AI RMF. AEGIS does not certify compliance and this report is not legal advice.

1. Overall maturity

Overall score	Maturity level	Interpretation
30.9 / 100	**Initial**	AI governance is ad hoc; immediate foundational work required.

2. Domain scores

Domain	Weight	Score	Profile
Strategy & Accountability	12%	33.3	[#######.....]
AI Inventory & Visibility	14%	26.7	[#####.....]
Risk Management & Lifecycle	14%	26.7	[#####.....]
Data & Privacy	12%	46.7	[#########.....]
Security & Resilience	14%	33.3	[#######.....]
Fairness, Safety & Transparency	12%	26.7	[#####.....]
Human Oversight & Operations	12%	33.3	[#######.....]
Assurance, Evidence & Improvement	10%	20.0	[####.....]

Scoring: domain score = sum of answers ÷ maximum possible × 100; overall = Σ(domain score × weight). Source: AEGIS self-assessment template AEGIS-SA-MATURITY v1.0.0.

3. Top strengths

- No practices were rated 4 or 5. Strengths will emerge as foundational controls are implemented.

4. Top gaps

#	Domain	Question	Answer	Recommended action	AEGIS controls
1	AI Inventory & Visibility	INV-Q1: Does your organization maintain a complete inventory of AI systems, models, agents, and vendor AI products?	1/5	Stand up an AI system of record and run inventory workshops plus shadow-AI discovery to reach full coverage.	EI-01, EI-08, EI-09
2	AI Inventory & Visibility	INV-Q2: Are AI systems classified by business criticality, data sensitivity, decision impact, autonomy, and regulatory relevance?	1/5	Apply AEGIS risk-tier criteria to every inventory record and approve tier assignments.	EI-02
3	Risk Management & Lifecycle	RSK-Q2: Are model, prompt, data, vendor, and agent changes subject to change control and reassessment?	1/5	Add AI-specific change triggers (model, prompt, data, vendor, tool, agent) to change management with reassessment rules.	GL-04, EI-05

#	Domain	Question	Answer	Recommended action	AEGIS controls
4	Risk Management & Lifecycle	RSK-Q3: Is there a defined process for retiring, restricting, or remediating AI systems that no longer meet requirements?	1/5	Define restrict/remediate/retire procedures and an exception register with expiry dates.	GL-09, GL-10
5	Security & Resilience	SEC-Q2: Are LLM applications and agents tested for prompt injection, jailbreak, data leakage, and unsafe tool use?	1/5	Run adversarial testing (prompt injection, jailbreak, leakage, unsafe tool use) before release and after material change.	IS-14, IS-11, IS-12

5. Critical risk flags

- **Critical gap — No reliable AI inventory** (trigger: INV-Q1). Without a complete inventory, risk tiering, oversight and evidence cannot be applied to systems you do not know about. Controls: EI-01, EI-09. Evidence status: User-asserted.
- **Critical gap — AI deployed without pre-deployment risk assessment** (trigger: RSK-Q1). Systems reaching production without risk assessment may carry unidentified legal, safety or security exposure. Controls: GL-02, GL-03. Evidence status: User-asserted.
- **Critical gap — Sensitive data may reach prompts, logs or external models** (trigger: DAT-Q3). Unredacted personal or confidential data in prompts and logs creates privacy and confidentiality exposure. Controls: IS-04, IS-03. Evidence status: User-asserted.
- **Critical gap — No tested kill switch or rollback** (trigger: SEC-Q3). Without containment, an AI incident cannot be stopped quickly. Controls: IS-15. Evidence status: User-asserted.
- **Critical gap — High-impact decisions lack defined human oversight** (trigger: HUM-Q1). High-impact AI decisions without review, override or appeal raise fairness, safety and accountability exposure. Controls: IS-09, IS-13. Evidence status: User-asserted.
- **Critical gap — Cannot produce audit-ready evidence** (trigger: ASR-Q2). Inability to evidence controls undermines responses to auditors, customers and regulators. Controls: SE-05, SE-08. Evidence status: User-asserted.
- **Critical gap — EU-facing AI without consistent AI-interaction disclosure** (trigger: FST-Q2). Transparency obligations for certain AI systems (EU AI Act Art. 50) require readiness; legal review recommended. Controls: IS-08. Evidence status: User-asserted.

6. 30/60/90-day action plan

First 30 days — stabilise

- Stand up an AI system of record and run inventory workshops plus shadow-AI discovery to reach full coverage. (AI Inventory & Visibility; INV-Q1; controls EI-01, EI-08, EI-09)
- Apply AEGIS risk-tier criteria to every inventory record and approve tier assignments. (AI Inventory & Visibility; INV-Q2; controls EI-02)
- Add AI-specific change triggers (model, prompt, data, vendor, tool, agent) to change management with reassessment rules. (Risk Management & Lifecycle; RSK-Q2; controls GL-04, EI-05)
- Define restrict/remediate/retire procedures and an exception register with expiry dates. (Risk Management & Lifecycle; RSK-Q3; controls GL-09, GL-10)
- Run adversarial testing (prompt injection, jailbreak, leakage, unsafe tool use) before release and after material change. (Security & Resilience; SEC-Q2; controls IS-14, IS-11, IS-12)

Days 31–60 — standardise

- Embed privacy impact assessment (consent, retention, transfers, minimisation) into AI intake and change. (Data & Privacy; DAT-Q2; controls IS-03)

Days 61–90 — scale and automate

- No actions in this horizon based on your answers.

7. Recommended AEGIS controls

Control	Pillar	Priority	NIST AI RMF	ISO/IEC 42001	EU AI Act
EI-01 Enterprise AI inventory (system of record)	EI	critical	GOVERN 1.6	A.4.2, 7.5	Art. 26, Art. 49
EI-09 Shadow AI discovery	EI	medium	GOVERN 1.6	A.4.2	—
GL-02 Pre-deployment risk and impact assessment	GL	critical	MAP 1.1, MAP 3.1, MAP 5.1, MEASURE 1.1	6.1.4, 8.4, A.5.2, A.5.3, A.5.4	Art. 9, Art. 27
GL-03 Risk-tiered approval gates	GL	critical	MANAGE 1.1, GOVERN 1.4	8.1, A.6.2.5	Art. 9, Art. 43
IS-04 Sensitive data redaction and masking	IS	critical	MEASURE 2.10	A.7.4, A.6.2.6	Art. 10
IS-03 Privacy assessment and data minimisation	IS	critical	MEASURE 2.10, MAP 1.1	A.7.4, A.5.4	Art. 10
IS-15 Kill switch and rollback	IS	critical	MANAGE 2.4	A.6.2.6	Art. 14
IS-09 Human oversight design	IS	critical	GOVERN 3.2, MAP 3.5	A.9.2, A.6.2.6	Art. 14, Art. 26
IS-13 Human approval for high-impact actions	IS	critical	GOVERN 3.2, MANAGE 2.2	A.9.2, A.6.2.6	Art. 14
SE-05 Audit-ready evidence packs	SE	high	GOVERN 1.4, MEASURE 2.9	7.5, 9.2	Art. 11, Art. 18

Framework mappings are indicative AEGIS interpretations (Control Library v1.0.0).

8. Framework readiness indicators (indicative — not a compliance determination)

NIST AI RMF

Area	Indicator	Band	Based on questions
GOVERN	28.0	Limited readiness	STR-Q1, STR-Q2, STR-Q3, HUM-Q3, ASR-Q3
MAP	32.0	Limited readiness	INV-Q1, INV-Q2, INV-Q3, RSK-Q1, DAT-Q1
MEASURE	32.0	Limited readiness	SEC-Q2, FST-Q1, HUM-Q2, ASR-Q1, DAT-Q2
MANAGE	28.0	Limited readiness	RSK-Q2, RSK-Q3, SEC-Q3, HUM-Q1, ASR-Q2

ISO/IEC 42001

Area	Indicator	Band	Based on questions
Context & leadership (Cl. 4–5, A.2–A.3)	33.3	Limited readiness	STR-Q1, STR-Q2, STR-Q3
Planning & impact assessment (Cl. 6, A.5)	26.7	Limited readiness	INV-Q2, RSK-Q1, FST-Q1
Support & documented information (Cl. 7, A.4)	20.0	Limited readiness	HUM-Q3, INV-Q1, ASR-Q2
Operation & lifecycle (Cl. 8, A.6–A.7)	40.0	Limited readiness	RSK-Q2, DAT-Q1, DAT-Q2, SEC-Q1
Performance evaluation (Cl. 9)	30.0	Limited readiness	HUM-Q2, ASR-Q1
Improvement (Cl. 10)	20.0	Limited readiness	ASR-Q3, RSK-Q3

EU AI Act

Area	Indicator	Band	Based on questions
AI literacy (Art. 4)	20.0	Limited readiness	HUM-Q3
Risk management (Art. 9)	26.7	Limited readiness	RSK-Q1, RSK-Q2, INV-Q2

Area	Indicator	Band	Based on questions
Data governance (Art. 10)	40.0	Limited readiness	DAT-Q1, DAT-Q2, FST-Q1
Technical documentation & record-keeping (Art. 11–12)	30.0	Limited readiness	INV-Q3, ASR-Q2
Transparency (Art. 13, Art. 50)	30.0	Limited readiness	FST-Q2, FST-Q3
Human oversight (Art. 14)	40.0	Limited readiness	HUM-Q1, SEC-Q3
Accuracy, robustness & cybersecurity (Art. 15)	30.0	Limited readiness	SEC-Q1, SEC-Q2
Post-market monitoring & incidents (Art. 72–73)	30.0	Limited readiness	HUM-Q2, ASR-Q1

These indicators summarise self-reported practices that relate to each framework area. They do not indicate conformity with ISO/IEC 42001, compliance with the EU AI Act, or alignment with NIST AI RMF, and EU AI Act applicability depends on your role (provider/deployer) and system classification.

9. Recommended AEGIS modules

- **AEGIS Assess** — Establish an evidence-based baseline of maturity and use-case risk.
- **AEGIS Register** — Create a single, owned and classified AI system of record.
- **AEGIS Control** — Operationalise policies, approval gates and lifecycle workflows.
- **AEGIS Shield** — Enforce runtime guardrails, approvals and containment for LLM apps and agents.
- **AEGIS Assurance** — Test controls continuously and produce audit-ready evidence packs.
- **AEGIS Agent** — Accelerate assessments, gap analysis and reporting with a citation-backed governance copilot.

10. Next step: AEGIS Governance Baseline

A fixed-scope 4–6 week engagement — inventory up to 25 AI systems, validate this maturity baseline with evidence, deep-dive five priority systems, review two agents, and deliver a gap register, remediation roadmap and executive readiness report.

Book an AEGIS Governance Baseline: /contact/book-governance-baseline/

Appendix A — Your answers

Question	Answer
STR-Q1: Does your organization have a documented AI governance policy and accountable ownership model?	2 — Initial / ad hoc
STR-Q2: Is there a cross-functional AI governance forum with authority to approve, restrict, or stop AI use cases?	1 — Not started
STR-Q3: Are AI roles, responsibilities, and escalation paths defined for business, technology, risk, legal, and security teams?	2 — Initial / ad hoc
INV-Q1: Does your organization maintain a complete inventory of AI systems, models, agents, and vendor AI products?	1 — Not started
INV-Q2: Are AI systems classified by business criticality, data sensitivity, decision impact, autonomy, and regulatory relevance?	1 — Not started
INV-Q3: Can you identify the owner, purpose, users, data sources, and dependencies for each material AI system?	2 — Initial / ad hoc
RSK-Q1: Do AI use cases undergo risk assessment before production deployment?	2 — Initial / ad hoc
RSK-Q2: Are model, prompt, data, vendor, and agent changes subject to change control and reassessment?	1 — Not started
RSK-Q3: Is there a defined process for retiring, restricting, or remediating AI systems that no longer meet requirements?	1 — Not started
DAT-Q1: Are data sources, data classifications, and data flows documented for AI systems?	2 — Initial / ad hoc
DAT-Q2: Are privacy, consent, retention, cross-border transfer, and data-minimization requirements assessed?	3 — Defined for some systems

Question	Answer
DAT-Q3: Are sensitive data fields redacted, masked, or excluded from prompts, logs, and external model calls?	2 — Initial / ad hoc
SEC-Q1: Are AI systems included in threat modeling, security review, and vulnerability-management processes?	2 — Initial / ad hoc
SEC-Q2: Are LLM applications and agents tested for prompt injection, jailbreak, data leakage, and unsafe tool use?	1 — Not started
SEC-Q3: Can you disable or roll back an AI system, model, agent, or tool quickly after a security or safety incident?	2 — Initial / ad hoc
FST-Q1: Are AI systems assessed for bias, discriminatory outcomes, safety risks, and unintended harm?	1 — Not started
FST-Q2: Are users informed when they are interacting with AI or when AI materially influences a decision?	2 — Initial / ad hoc
FST-Q3: Are explanations, disclosures, and appeal mechanisms available for affected individuals where appropriate?	1 — Not started
HUM-Q1: Are human-review, override, escalation, and appeal mechanisms defined for high-impact AI decisions?	2 — Initial / ad hoc
HUM-Q2: Are AI systems monitored for performance, drift, misuse, incidents, and unexpected behavior?	2 — Initial / ad hoc
HUM-Q3: Are operators and affected teams trained on AI limitations, escalation paths, and responsible use?	1 — Not started
ASR-Q1: Are AI controls tested periodically, with evidence retained for audit and review?	1 — Not started
ASR-Q2: Can your organization produce an audit-ready evidence pack for a material AI system on request?	1 — Not started
ASR-Q3: Are AI governance metrics reviewed by leadership and used to drive corrective action?	1 — Not started

Appendix B — Sources

- AEGIS Control EI-01: Enterprise AI inventory (system of record) (Control Library v1.0.0)
- AEGIS Control EI-09: Shadow AI discovery (Control Library v1.0.0)
- AEGIS Control GL-02: Pre-deployment risk and impact assessment (Control Library v1.0.0)
- AEGIS Control GL-03: Risk-tiered approval gates (Control Library v1.0.0)
- AEGIS Control IS-04: Sensitive data redaction and masking (Control Library v1.0.0)
- AEGIS Control IS-03: Privacy assessment and data minimisation (Control Library v1.0.0)
- AEGIS Control IS-15: Kill switch and rollback (Control Library v1.0.0)
- AEGIS Control IS-09: Human oversight design (Control Library v1.0.0)
- AEGIS Control IS-13: Human approval for high-impact actions (Control Library v1.0.0)
- AEGIS Control SE-05: Audit-ready evidence packs (Control Library v1.0.0)
- AEGIS Control SE-08: Immutable action logging (Control Library v1.0.0)
- AEGIS Control IS-08: Transparency and AI disclosures (Control Library v1.0.0)
- AEGIS Self-Assessment template AEGIS-SA-MATURITY v1.0.0

AEGIS supports governance readiness, control implementation and evidence collection. It does not certify ISO/IEC 42001 conformity, guarantee EU AI Act compliance, or replace legal advice. Framework mappings are indicative.