Executive Readiness Report — Example Regional Bank (synthetic)

> ***SYNTHETIC DRESS REHEARSAL.*** *Fictional organization, generated answers. This report demonstrates the AEGIS Pilot A workflow and outputs; it is not evidence about any real organization.*

Prepared: 2026-10-21 · **Scope:** AEGIS Governance Baseline (Pilot A profile) · **Prepared by:** AEGIS Agent-assisted workflow, synthetic analyst

1. Executive summary

- **28 AI systems** registered as draft inventory records (EI-01); preliminary tiers: High 12, Medium 14, Low 2. Tiers are Inferred and were confirmed through the approval queue (EI-02, IS-13).
- **Enterprise maturity: 48.1 / 100 — Developing** (self-reported, User-asserted).
- **12 deep-dive assessments** covering every High-tier system; **403 gaps** identified (Critical 40, High 254, Medium 90, Low 19).
- **2 runtime policies** drafted for agentic systems; both remain in draft/observe mode pending human approval to activate (IS-12, IS-13).
- Evidence position: Missing evidence: 262, User-asserted: 101, Critical gap: 40.

2. AI portfolio by risk tier

Tier	Systems	Examples
High	12	SME credit decisioning model, Retail credit scorecard, Card fraud detection, KYC document verification
Medium	14	AML transaction monitoring, Chargeback dispute assistant, Marketing propensity model, Internal code assistant
Low	2	Meeting summariser, Branch staffing forecast

3. Maturity baseline by domain

Domain	Score
Strategy & Accountability	53.3
AI Inventory & Visibility	40.0
Risk Management & Lifecycle	46.7
Data & Privacy	53.3
Security & Resilience	46.7
Fairness, Safety & Transparency	46.7
Human Oversight & Operations	53.3
Assurance, Evidence & Improvement	46.7

4. Risk heat map — gaps by system and severity

System	Tier	Critical	High	Medium	Low
Retail credit scorecard	High	5	21	5	2
KYC document verification	High	4	19	5	2
Biometric face match (ID&V)	High	2	25	3	1
Collections outreach agent	High	6	26	8	2
CV screening assistant	High	4	23	9	2
SME credit decisioning model	High	1	19	8	1
Card fraud detection	High	2	24	7	3

System	Tier	Critical	High	Medium	Low
Customer support agent	High	6	25	10	1
Mortgage valuation model	High	1	22	8	1
Robo-advice portfolio tool	High	4	17	8	2
Employee attrition model	High	3	17	7	1
Cyber threat detection	High	2	16	12	1

5. Most frequent Critical/High control gaps

Control	Occurrences
IS-03 Privacy assessment and data minimisation	28
EI-05 Model, prompt and component registry	25
SE-04 AI incident management and reporting	19
SE-02 Evidence repository and retention	18
IS-08 Transparency and AI disclosures	14
GL-11 Vendor AI due diligence and contract controls	14
AO-09 Third-party AI accountability	14
GL-04 Change management and reassessment triggers	13
IS-06 Safety and harm evaluation	13
IS-10 Secure model and prompt supply chain	12

6. Remediation roadmap

- **First 30 days:** close Critical gaps (owners and due dates in the gap register), stand up the AI system of record as authoritative (EI-01, EI-03), enable redaction and kill switch for agentic systems (IS-04, IS-15).
- **Days 31–60:** High gaps; adversarial testing of LLM apps and agents (IS-14); change triggers for model/prompt/tool changes (GL-04); move runtime policies from observe to warn/protect (IS-12).
- **Days 61–90:** Medium gaps; control-testing programme and evidence repository (SE-01, SE-02); leadership KPI pack (SE-03); enforce mode for agents after approval (IS-13).

7. Runtime Shield recommendations

Agent	Policy	Version	Status	Rollout mode	Risk rating	Activation
Customer support agent	pol-customer-support-agent	0.1.0	draft	observe	High	apr_mvi8hwkv3b5b0409 (pending)
Collections outreach agent	pol-collections-outreach-agent	0.1.0	draft	observe	Critical	apr_mvi8hwlj3ae88ab5 (pending)

8. Pilot success metrics (dry run)

Metric	Target	Result	Status
AI systems registered	≥ 25	28	Met (synthetic)
High-risk systems assessed	100 %	12/12 (100%)	Met (synthetic)
Agent-assisted assessment completion	≥ 70%	100% (all responses captured via agent tools; scripted)	Met (synthetic — not a measure of user adoption)
Critical gaps documented with owners	Yes	40 critical gaps, all with owner role + due date	Met (synthetic)
Runtime policies drafted	≥ 2	2	Met (synthetic)

Metric	Target	Result	Status
Executive report accepted	Yes	Report generated	Requires real partner
Follow-on commercial discussion initiated	Yes	—	Requires real partner

9. Governance checks exercised

- W2: registered 28 AI systems as draft records
- W2: 28 tier confirmations approved by approver; separation-of-duties check: PASS — self-approval refused (ApprovalError)
- W3: maturity baseline 48.1 (Developing), 6 critical flags
- W4: completed 12 deep-dive assessments via agent tools (419 responses)
- W4: 2 agent reviews; runtime policies drafted (status draft, activation pending human approval)
- W5: gap register 403 gaps; 40 remediation tasks created for Critical gaps after explicit confirmation
- W6: evidence pack export: first call -> pending_approval; after approval -> exported (7 items)

10. Artefacts

- `01_ai_inventory.csv` — inventory with preliminary tiers
- `02_maturity_baseline.{md,pdf}`
- `03_deep_dives/<system>/` — per-system report (md, pdf, csv gap register, json)
- `04_runtime_policies/` — advisor output + policy YAML/JSON for two agents
- `05_gap_register.csv` — consolidated
- `06_evidence_pack/` — approval-gated evidence pack
- `metrics.json`, `run_log.md`

AEGIS supports governance readiness, control implementation and evidence collection. It does not certify ISO/IEC 42001 conformity, guarantee EU AI Act compliance, or replace legal advice. Framework mappings are indicative.