AI Use-Case Governance Readiness Report — Card fraud detection

Synthetic example. *This report was generated from fictional data to illustrate AEGIS Agent output. It describes no real organization or system.*

Field	Value
Report ID	rpt_mvi8hui0a98a2ef6
Assessment ID	asm_mvi8hudz8f5cf5f8
AI system record	sys_mvi8hs7he2b521eb
Generated (UTC)	2026-10-21T14:55:23+00:00
Prepared by	Synthetic AEGIS Analyst via AEGIS Agent
Question bank	AEGIS-UC-ASSESS v1.0.0
Control library	AEGIS Control Library v1.0.0
Evidence basis	Self-reported answers (User-asserted) unless evidence was reviewed by a named reviewer (Verified)

1. Executive summary

- **Preliminary risk tier: High** (Inferred; requires human confirmation — Confirmed by Synthetic Head of AI Risk (2026-10-21)).
- **38 of 38** applicable questions answered across 8 modules.
- **36 gaps** identified: 2 Critical, 24 High, 7 Medium, 3 Low.
- **Evidence posture:** 2 finding(s) Verified by a named reviewer; 10 User-asserted; 24 Missing evidence; 2 Critical gap; 0 Inferred.
- **Priority actions (Critical, due within 30 days):**
  - GAP-001: Require human approval for high-impact actions via runtime policy, and record approver identity for each approval. (IS-13; owner: Business Owner)
  - GAP-002: Document the oversight design (reviewer roles, review points, override and stop authority) and keep records of oversight in operation. (IS-09; owner: Business Owner)

This report indicates governance readiness and evidence gaps for one AI system. It is not a determination of conformity or legal compliance with any framework or law.

2. AI system profile

Field	Value	Evidence status
name	Card fraud detection	User-asserted
purpose	Flag suspicious card transactions for review	User-asserted
domain_category	financial_services	User-asserted
decision_impact	material	User-asserted
affected_parties	customers	User-asserted
data_sensitivity	personal	User-asserted
autonomy	acts_with_approval	User-asserted
external_actions	No	User-asserted
genai	No	User-asserted
vendor_model	Yes	User-asserted
eu_exposure	Yes	User-asserted
tools	card_block	User-asserted

Field	Value	Evidence status
environment	production	User-asserted

3. Risk tier and rationale

Tier: High (score 85) — evidence status: Inferred; requires human confirmation: yes.

Rationale: Preliminary tier High (score 85). Drivers: decision_impact=material (+20); affected_parties=customers (+10); data_sensitivity=personal (+15); autonomy=acts_with_approval (+15); tool_access=1 tool(s), 0 with write/act capability (+5); domain_category=financial_services (+10); third_party_model=True (+5); eu_exposure=True (+5). Inferred from user-provided intake; requires human confirmation.

Factor	Value	Points	Citation
decision_impact	material	20	AEGIS Control GL-02: Pre-deployment risk and impact assessment (Control Library v1.0.0)
affected_parties	customers	10	AEGIS Control EI-04: Intended purpose and context of use (Control Library v1.0.0)
data_sensitivity	personal	15	AEGIS Control IS-03: Privacy assessment and data minimisation (Control Library v1.0.0)
autonomy	acts_with_approval	15	AEGIS Control IS-09: Human oversight design (Control Library v1.0.0)
tool_access	1 tool(s), 0 with write/act capability	5	AEGIS Control IS-11: Least-privilege tool access (Control Library v1.0.0)
domain_category	financial_services	10	AEGIS Control AO-07: Regulatory and contractual obligations register (Control Library v1.0.0)
third_party_model	Yes	5	AEGIS Control GL-11: Vendor AI due diligence and contract controls (Control Library v1.0.0)
eu_exposure	Yes	5	AEGIS Control AO-07: Regulatory and contractual obligations register (Control Library v1.0.0)

- Approval authority for this tier: AI Governance Forum
- Review cadence: every 6 months
- Assessment modules applied: Core governance & lifecycle, Data & privacy, Security & resilience, Agentic AI & tool use, Fairness, safety & transparency, Human oversight & operations, Vendor & third-party AI, Monitoring, incidents & evidence
- Method: AEGIS risk-tier rules v1.0 (thresholds: High >= 60 or any high-risk trigger; Medium >= 30)

4. Domain heat map

Domain	Qs	Implementation (self-reported)	Verified	Crit	High	Med	Low	Rating
Strategy & Accountability	2	50.0%	0	0	1	1	0	High exposure
AI Inventory & Visibility	4	25.0%	0	0	3	1	0	High exposure
Risk Management & Lifecycle	7	35.7%	0	0	5	1	1	High exposure
Data & Privacy	5	60.0%	0	0	3	2	0	High exposure
Security & Resilience	6	66.7%	1	0	2	1	2	High exposure
Fairness, Safety & Transparency	4	37.5%	0	0	3	1	0	High exposure
Human Oversight & Operations	6	25.0%	1	2	3	0	0	Critical exposure
Assurance, Evidence & Improvement	4	25.0%	0	0	4	0	0	High exposure

Implementation indicator = implemented (1.0) and partial (0.5) answers as a share of applicable questions; it reflects self-reported answers, not verified evidence.

5. Gap register

Gap · severity · due	Evidence status	Question	Controls	Owner role	Recommended action
GAP-001 · Critical · due 2026-11-20	Critical gap	UC-AGT-Q3: Do high-impact actions (e.g. refunds, record changes, external messages, deletions) require human approval before execution?	IS-13	Business Owner	Require human approval for high-impact actions via runtime policy, and record approver identity for each approval.
GAP-002 · Critical · due 2026-11-20	Critical gap	UC-HUM-Q1: Is the human oversight design documented — who reviews which outputs or actions, when, and with authority to override or stop the system?	IS-09	Business Owner	Document the oversight design (reviewer roles, review points, override and stop authority) and keep records of oversight in operation.
GAP-003 · High · due 2026-12-20	Missing evidence	UC-CORE-Q1: Is there a named accountable business owner and technical owner for this AI system, recorded in the AI inventory?	EI-03, EI-01	Business unit leaders	Register the system in the AI inventory with a named accountable business owner and technical owner, and capture owner acceptance.
GAP-004 · High · due 2026-12-20	Missing evidence	UC-CORE-Q4: Was a pre-deployment risk and impact assessment completed and approved at the gate appropriate to the risk tier?	GL-02, GL-03	AI Risk Lead	Complete a tier-proportionate risk and impact assessment and obtain approval at the tier-appropriate gate before (further) production use.
GAP-005 · High · due 2026-12-20	Missing evidence	UC-DATA-Q2: Has a privacy / data-protection impact assessment been completed covering lawful basis, minimisation, retention and cross-border transfers?	IS-03	DPO / Privacy Lead	Complete a privacy impact assessment (DPIA where required) covering lawful basis, minimisation, retention and transfers, signed off by the DPO/privacy lead.
GAP-006 · High · due 2026-12-20	Missing evidence	UC-DATA-Q3: Are personal and sensitive fields redacted or masked before they reach prompts, logs and external model calls?	IS-04	Platform Engineering / Privacy Lead	Deploy redaction/masking on prompts, responses, logs and external model calls; test it against representative records and retain the results.
GAP-007 · High · due 2026-12-20	Missing evidence	UC-DATA-Q4: Is the system's data access limited to the minimum records and fields it needs (e.g. field-level scoping of CRM access)?	IS-03, IS-11	DPO / Privacy Lead	Restrict data access to the minimum records and fields required, using scoped service identities, and review access periodically.
GAP-008 · High · due 2026-12-20	Missing evidence	UC-AGT-Q4: Are all agent tool calls logged in a tamper-evident log with trace IDs, redacted inputs and outcomes?	SE-08	Platform Engineering / Security Operations	Log every tool call with trace ID, redacted inputs, decision and outcome to tamper-evident storage with defined retention.
GAP-009 · High · due 2026-12-20	Missing evidence	UC-CORE-Q2: Are the intended purpose, context of use, users, affected parties and out-of-scope uses documented?	EI-04	Product Owner	Document intended purpose, context of use, users, affected parties and explicitly out-of-scope or foreseeable misuse.
GAP-010 · High · due 2026-12-20	Missing evidence	UC-CORE-Q5: Have the legal, regulatory and contractual obligations that apply to this system (by jurisdiction) been identified and reviewed?	AO-07	Legal / Compliance	Add the system to the obligations register and obtain a legal review of applicable obligations by jurisdiction.
GAP-011 · High · due 2026-12-20	Missing evidence	UC-CORE-Q7: Are changes to the model, prompts, data, tools or vendor subject to change control with defined reassessment triggers?	GL-04, EI-05	Engineering / Change Advisory Board	Add AI-specific change triggers (model, prompt, data, tool, vendor) to change management and require reassessment when they fire.
GAP-012 · High · due 2026-12-20	Missing evidence	UC-SEC-Q1: Has an AI-specific threat model been completed for this system (e.g. prompt injection, data poisoning, model abuse, tool misuse)?	IS-01	CISO / Security Architecture	Produce an AI-specific threat model covering injection, poisoning, abuse and tool misuse, and track mitigations to closure.

Gap · severity · due	Evidence status	Question	Controls	Owner role	Recommended action
GAP-013 · High · due 2026-12-20	Missing evidence	UC-SEC-Q2: Do models, prompts, plugins and third-party components come from an approved, integrity-checked supply chain?	IS-10, EI-05	Security Architecture	Restrict models and components to an approved list with integrity checks and dependency scanning.
GAP-014 · High · due 2026-12-20	Missing evidence	UC-FAIR-Q1: Has the system been tested for bias or discriminatory outcomes across relevant groups, with documented thresholds?	IS-05	Responsible AI Lead / Data Science	Define relevant groups and fairness metrics, test against thresholds and record mitigations.
GAP-015 · High · due 2026-12-20	Missing evidence	UC-FAIR-Q2: Has a safety and harm evaluation been completed (harmful, misleading or inappropriate outputs, including for vulnerable users)?	IS-06	Responsible AI Lead	Complete a harm analysis and safety evaluation, including vulnerable-user scenarios, and obtain residual-risk sign-off.
GAP-016 · High · due 2026-12-20	Missing evidence	UC-FAIR-Q3: Are people informed when they are interacting with AI or receiving AI-generated content?	IS-08	Product Owner / Legal	Add clear AI-interaction and AI-generated-content disclosures at the point of interaction; have Legal review the wording.
GAP-017 · High · due 2026-12-20	Missing evidence	UC-HUM-Q2: Are operators and reviewers trained on the system's limitations, automation bias and escalation paths?	AO-05	HR / Learning lead with AI Governance Lead	Deliver role-based training for operators and reviewers covering limitations, automation bias and escalation, and track completion.
GAP-018 · High · due 2026-12-20	Missing evidence	UC-HUM-Q3: Is there a defined escalation path for concerns and AI incidents involving this system?	AO-06, SE-04	AI Governance Lead	Define and communicate an escalation path and incident procedure for this system, with named contacts and timeframes.
GAP-019 · High · due 2026-12-20	Missing evidence	UC-HUM-Q4: Were go-live deployment conditions defined, verified and signed off before production use?	GL-07	Product Owner	Define go-live conditions (controls, monitoring, oversight) and verify them with a signed checklist before production use.
GAP-020 · High · due 2026-12-20	Missing evidence	UC-VEN-Q2: Do contract terms cover data-use restrictions, no training on your data, retention, breach notification and audit rights?	GL-11, AO-09	Procurement / Vendor Risk	Negotiate or confirm AI contract clauses for data use, training restrictions, retention, breach notification and audit rights.
GAP-021 · High · due 2026-12-20	Missing evidence	UC-VEN-Q3: Is the vendor AI product recorded in the third-party AI inventory with an accountable owner?	EI-08, AO-09	Procurement / Vendor Risk	Record the vendor AI product in the third-party AI inventory with an accountable owner and contract/DPA references.
GAP-022 · High · due 2026-12-20	Missing evidence	UC-VEN-Q4: Are vendor model version changes and deprecations monitored and routed through change control?	GL-04, EI-05	Engineering / Change Advisory Board	Monitor vendor model/version announcements and route changes through change control with reassessment triggers.
GAP-023 · High · due 2026-12-20	Missing evidence	UC-MON-Q2: Is this system covered by an AI incident management procedure, including internal and external reporting timelines?	SE-04	CISO / Operational Risk	Bring the system under the AI incident procedure with defined severity levels and reporting timelines; Legal to confirm external reporting duties.
GAP-024 · High · due 2026-12-20	Missing evidence	UC-MON-Q3: Is evidence for this system's controls retained in an evidence repository, and can an audit-ready evidence pack be produced on request?	SE-02, SE-05	AI Governance Lead	Store control evidence for this system in the evidence repository with retention rules and assemble an audit-ready evidence pack.
GAP-025 · High · due 2026-12-20	Missing evidence	UC-MON-Q5: Is technical documentation (system/model card, architecture, data, testing, limitations) maintained for this system?	GL-06	Product / ML Engineering	Produce and maintain technical documentation including a system/model card, architecture, data, testing and known limitations.

Gap · severity · due	Evidence status	Question	Controls	Owner role	Recommended action
GAP-026 · High · due 2026-12-20	Missing evidence	UC-MON-Q4: Is a periodic review / recertification of this system scheduled at the cadence required by its risk tier?	GL-12, EI-11	AI Risk Lead	Schedule periodic review and recertification at the cadence set for the system's risk tier.
GAP-027 · Medium · due 2027-01-19	User-asserted	UC-CORE-Q3: Has the use case passed intake triage and received a risk tier that was confirmed by an accountable approver?	GL-01, EI-02, EI-12	AI Risk Lead	Run the use case through intake triage and record a human-confirmed risk tier with rationale in the intake register.
GAP-028 · Medium · due 2027-01-19	User-asserted	UC-CORE-Q6: Has the use case been screened against the organization's prohibited and restricted AI use policy?	AO-10	Legal / AI Governance Lead	Screen the use case against the prohibited/restricted AI use list and record the outcome; escalate any match for legal review.
GAP-029 · Medium · due 2027-01-19	User-asserted	UC-DATA-Q5: Are retention periods defined and enforced for prompts, outputs, conversation transcripts and logs?	IS-03, SE-02	DPO / Privacy Lead	Define and enforce retention periods for prompts, outputs, transcripts and logs consistent with the privacy assessment.
GAP-030 · Medium · due 2027-01-19	User-asserted	UC-AGT-Q1: Is the agent registered with every tool it can call, the scopes of each tool and the permitted actions?	EI-07	Platform Engineering Lead	Register the agent and a tool/permission manifest listing each tool, scope and permitted action.
GAP-031 · Medium · due 2027-01-19	User-asserted	UC-AGT-Q2: Does each tool use least-privilege credentials scoped to the required actions only (no shared or admin credentials)?	IS-11	Platform Engineering / IAM	Replace shared/admin credentials with per-tool, least-privilege scoped identities and review access periodically.
GAP-032 · Medium · due 2027-01-19	User-asserted	UC-DATA-Q1: Are the data sources, classifications and data flows used for training, grounding (RAG) and prompts documented for this system?	EI-06, IS-02	Data Owner / CDO office	Document every data source, its classification and the data flows into prompts, retrieval indexes, logs and external services.
GAP-033 · Medium · due 2027-01-19	User-asserted	UC-FAIR-Q4: Can affected individuals obtain an explanation, reach a human, and contest or appeal outcomes?	IS-08, AO-11	Product Owner / Legal	Provide a route to a human, explanations proportionate to impact, and a contest/appeal channel whose outcomes are logged.
GAP-034 · Low · due 2027-04-19	User-asserted	UC-SEC-Q4: Were pre-release validation and robustness tests run against documented acceptance criteria, with results retained?	GL-05, IS-07	QA / ML Engineering	Define acceptance criteria, run pre-release validation and robustness tests, and retain signed-off results.
GAP-035 · Low · due 2027-04-19	User-asserted	UC-AGT-Q5: Are per-action limits (value caps, rate limits, loop/step limits) enforced at runtime for agent actions?	IS-12, IS-16	Platform Engineering / Security	Enforce value caps, rate limits and step limits for agent actions in the runtime policy with alerts on breach.
GAP-036 · Low · due 2027-04-19	User-asserted	UC-VEN-Q1: Was vendor AI due diligence completed for the model/service provider (security, data use, model training on your data, subprocessors)?	GL-11	Procurement / Vendor Risk	Complete AI-specific vendor due diligence covering security, data use, training on customer data and subprocessors.

Due dates apply the AEGIS severity SLAs (Critical 30, High 60, Medium 90, Low 180 days) from 2026-10-21.

6. Control recommendations

Control	Title	Priority	Highest gap severity	Owner role	Framework mappings (indicative)
IS-09	Human oversight design	critical	Critical	Business Owner	NIST AI RMF: GOVERN 3.2, MAP 3.5; ISO/IEC 42001: A.9.2, A.6.2.6; EU AI Act: Art. 14, Art. 26

Control	Title	Priority	Highest gap severity	Owner role	Framework mappings (indicative)
IS-13	Human approval for high-impact actions	critical	Critical	Business Owner	NIST AI RMF: GOVERN 3.2, MANAGE 2.2; ISO/IEC 42001: A.9.2, A.6.2.6; EU AI Act: Art. 14
EI-01	Enterprise AI inventory (system of record)	critical	High	AI Governance Lead / CIO office	NIST AI RMF: GOVERN 1.6; ISO/IEC 42001: A.4.2, 7.5; EU AI Act: Art. 26, Art. 49
EI-03	Accountable owner per AI system	critical	High	Business unit leaders	NIST AI RMF: GOVERN 2.1; ISO/IEC 42001: A.3.2, A.10.2; EU AI Act: Art. 26
GL-02	Pre-deployment risk and impact assessment	critical	High	AI Risk Lead	NIST AI RMF: MAP 1.1, MAP 3.1, MAP 5.1, MEASURE 1.1; ISO/IEC 42001: 6.1.4, 8.4, A.5.2, A.5.3, A.5.4; EU AI Act: Art. 9, Art. 27
GL-03	Risk-tiered approval gates	critical	High	AI Governance Forum	NIST AI RMF: MANAGE 1.1, GOVERN 1.4; ISO/IEC 42001: 8.1, A.6.2.5; EU AI Act: Art. 9, Art. 43
IS-03	Privacy assessment and data minimisation	critical	High	DPO / Privacy Lead	NIST AI RMF: MEASURE 2.10, MAP 1.1; ISO/IEC 42001: A.7.4, A.5.4; EU AI Act: Art. 10
IS-04	Sensitive data redaction and masking	critical	High	Platform Engineering / Privacy Lead	NIST AI RMF: MEASURE 2.10; ISO/IEC 42001: A.7.4, A.6.2.6; EU AI Act: Art. 10
IS-11	Least-privilege tool access	critical	High	Platform Engineering / IAM	NIST AI RMF: MEASURE 2.7, MANAGE 2.2; ISO/IEC 42001: A.6.2.6, A.4.4; EU AI Act: Art. 15, Art. 14
SE-08	Immutable action logging	critical	High	Platform Engineering / Security Operations	NIST AI RMF: MEASURE 2.8, MANAGE 4.1; ISO/IEC 42001: A.6.2.8; EU AI Act: Art. 12, Art. 19, Art. 26
AO-05	AI literacy and role-based training	high	High	HR / Learning lead with AI Governance Lead	NIST AI RMF: GOVERN 2.2; ISO/IEC 42001: 7.2, 7.3, A.4.6; EU AI Act: Art. 4
AO-06	Escalation paths and decision rights	high	High	AI Governance Lead	NIST AI RMF: GOVERN 2.1, GOVERN 4.3; ISO/IEC 42001: 5.3, A.3.3; EU AI Act: Art. 26
AO-07	Regulatory and contractual obligations register	high	High	Legal / Compliance	NIST AI RMF: GOVERN 1.1; ISO/IEC 42001: 4.1, 4.2, 6.1; EU AI Act: Art. 6, Art. 26
EI-04	Intended purpose and context of use	high	High	Product Owner	NIST AI RMF: MAP 1.1, MAP 1.2, MAP 3.4; ISO/IEC 42001: A.6.2.2, A.9.4; EU AI Act: Art. 11, Art. 13
EI-05	Model, prompt and component registry	high	High	ML / Platform Engineering Lead	NIST AI RMF: MAP 4.1, GOVERN 1.6; ISO/IEC 42001: A.4.4, A.6.2.3; EU AI Act: Art. 11
EI-08	Third-party and vendor AI inventory	high	High	Procurement / Vendor Risk	NIST AI RMF: GOVERN 6.1; ISO/IEC 42001: A.10.3; EU AI Act: Art. 25
GL-04	Change management and reassessment triggers	high	High	Engineering / Change Advisory Board	NIST AI RMF: GOVERN 1.4, MANAGE 2.4; ISO/IEC 42001: 6.3, A.6.2.6; EU AI Act: Art. 43
GL-06	Technical documentation	high	High	Product / ML Engineering	NIST AI RMF: MAP 2.2, MEASURE 2.9; ISO/IEC 42001: 7.5, A.6.2.7, A.6.2.3; EU AI Act: Art. 11
GL-07	Deployment conditions and go-live readiness	high	High	Product Owner	NIST AI RMF: MANAGE 1.2, MANAGE 2.1; ISO/IEC 42001: A.6.2.5; EU AI Act: Art. 26
GL-11	Vendor AI due diligence and contract controls	high	High	Procurement / Vendor Risk	NIST AI RMF: GOVERN 6.1, MAP 4.1; ISO/IEC 42001: A.10.3, A.10.4; EU AI Act: Art. 25
IS-01	AI threat modelling	high	High	CISO / Security Architecture	NIST AI RMF: MAP 5.1, MEASURE 2.7; ISO/IEC 42001: 6.1.2, A.6.1.3; EU AI Act: Art. 15

Control	Title	Priority	Highest gap severity	Owner role	Framework mappings (indicative)
IS-05	Bias and fairness testing	high	High	Responsible AI Lead / Data Science	NIST AI RMF: MEASURE 2.11; ISO/IEC 42001: A.5.4, A.7.4; EU AI Act: Art. 10
IS-06	Safety and harm evaluation	high	High	Responsible AI Lead	NIST AI RMF: MEASURE 2.6, MAP 5.1; ISO/IEC 42001: A.5.4, A.5.5; EU AI Act: Art. 9
IS-08	Transparency and AI disclosures	high	High	Product Owner / Legal	NIST AI RMF: MEASURE 2.8, MEASURE 2.9; ISO/IEC 42001: A.8.2, A.8.5; EU AI Act: Art. 13, Art. 50
IS-10	Secure model and prompt supply chain	high	High	Security Architecture	NIST AI RMF: GOVERN 6.1, MEASURE 2.7; ISO/IEC 42001: A.10.3, A.6.2.3; EU AI Act: Art. 15
SE-02	Evidence repository and retention	high	High	AI Governance Lead	NIST AI RMF: GOVERN 1.4; ISO/IEC 42001: 7.5; EU AI Act: Art. 18, Art. 19
SE-04	AI incident management and reporting	high	High	CISO / Operational Risk	NIST AI RMF: MANAGE 4.3, MANAGE 2.3; ISO/IEC 42001: 10.2, A.8.4; EU AI Act: Art. 73, Art. 26
SE-05	Audit-ready evidence packs	high	High	AI Governance Lead	NIST AI RMF: GOVERN 1.4, MEASURE 2.9; ISO/IEC 42001: 7.5, 9.2; EU AI Act: Art. 11, Art. 18
AO-09	Third-party AI accountability	medium	High	Procurement / Vendor Risk	NIST AI RMF: GOVERN 6.1, GOVERN 6.2; ISO/IEC 42001: A.10.2, A.10.3; EU AI Act: Art. 25, Art. 26
EI-11	Inventory attestation and recertification	medium	High	AI Governance Lead	NIST AI RMF: GOVERN 1.5; ISO/IEC 42001: 9.1
GL-12	Periodic review and recertification	medium	High	AI Risk Lead	NIST AI RMF: GOVERN 1.5; ISO/IEC 42001: 9.3; EU AI Act: Art. 9
AO-10	Prohibited and restricted AI use policy	critical	Medium	Legal / AI Governance Lead	NIST AI RMF: GOVERN 1.1, MAP 1.1; ISO/IEC 42001: A.9.4, A.2.2; EU AI Act: Art. 5
EI-02	AI risk classification and tiering	critical	Medium	AI Risk Lead	NIST AI RMF: MAP 1.1, MAP 1.5, GOVERN 1.3; ISO/IEC 42001: 6.1.2, A.5.2; EU AI Act: Art. 6
EI-07	Agent and tool registry	critical	Medium	Platform Engineering Lead	NIST AI RMF: GOVERN 1.6, MAP 4.1; ISO/IEC 42001: A.4.4, A.6.2.3; EU AI Act: Art. 11
EI-06	Data source and lineage registry	high	Medium	Data Owner / CDO office	NIST AI RMF: MAP 2.3, MEASURE 2.10; ISO/IEC 42001: A.4.3, A.7.5; EU AI Act: Art. 10
EI-12	Use-case intake register	high	Medium	AI CoE	NIST AI RMF: MAP 1.1; ISO/IEC 42001: A.5.2, A.6.1.3; EU AI Act: Art. 6
GL-01	Intake triage workflow	high	Medium	AI CoE	NIST AI RMF: MAP 1.1, MAP 1.6; ISO/IEC 42001: 8.1, A.6.1.3; EU AI Act: Art. 9
IS-02	Data governance and quality	high	Medium	Data Owner	NIST AI RMF: MAP 2.3, MEASURE 2.1; ISO/IEC 42001: A.7.2, A.7.3, A.7.4, A.7.6; EU AI Act: Art. 10
AO-11	Stakeholder engagement and feedback	medium	Medium	Product Owner	NIST AI RMF: GOVERN 5.1, GOVERN 5.2; ISO/IEC 42001: 4.2, A.8.3, A.8.5; EU AI Act: Art. 27
GL-05	Pre-release validation and testing	high	Low	QA / ML Engineering	NIST AI RMF: MEASURE 2.1, MEASURE 2.3, MEASURE 2.5; ISO/IEC 42001: A.6.2.4; EU AI Act: Art. 9, Art. 15
IS-07	Accuracy and robustness testing	high	Low	ML Engineering	NIST AI RMF: MEASURE 2.5; ISO/IEC 42001: A.6.2.4; EU AI Act: Art. 15
IS-12	Runtime guardrails and policy enforcement	high	Low	Platform Engineering / Security	NIST AI RMF: MANAGE 2.2, MANAGE 2.4, MEASURE 2.7; ISO/IEC 42001: A.6.2.6; EU AI Act: Art. 15, Art. 14

Control	Title	Priority	Highest gap severity	Owner role	Framework mappings (indicative)
IS-16	Rate, budget and resource limits	medium	Low	Platform Engineering / FinOps	NIST AI RMF: MANAGE 2.2; ISO/IEC 42001: A.4.5, A.6.2.6; EU AI Act: Art. 15

7. Remediation roadmap (30/60/90 days)

First 30 days — contain Critical exposure

- GAP-001 (Critical): Require human approval for high-impact actions via runtime policy, and record approver identity for each approval. — IS-13; owner: Business Owner; due 2026-11-20
- GAP-002 (Critical): Document the oversight design (reviewer roles, review points, override and stop authority) and keep records of oversight in operation. — IS-09; owner: Business Owner; due 2026-11-20

Days 31–60 — close High gaps

- GAP-003 (High): Register the system in the AI inventory with a named accountable business owner and technical owner, and capture owner acceptance. — EI-03, EI-01; owner: Business unit leaders; due 2026-12-20
- GAP-004 (High): Complete a tier-proportionate risk and impact assessment and obtain approval at the tier-appropriate gate before (further) production use. — GL-02, GL-03; owner: AI Risk Lead; due 2026-12-20
- GAP-005 (High): Complete a privacy impact assessment (DPIA where required) covering lawful basis, minimisation, retention and transfers, signed off by the DPO/privacy lead. — IS-03; owner: DPO / Privacy Lead; due 2026-12-20
- GAP-006 (High): Deploy redaction/masking on prompts, responses, logs and external model calls; test it against representative records and retain the results. — IS-04; owner: Platform Engineering / Privacy Lead; due 2026-12-20
- GAP-007 (High): Restrict data access to the minimum records and fields required, using scoped service identities, and review access periodically. — IS-03, IS-11; owner: DPO / Privacy Lead; due 2026-12-20
- GAP-008 (High): Log every tool call with trace ID, redacted inputs, decision and outcome to tamper-evident storage with defined retention. — SE-08; owner: Platform Engineering / Security Operations; due 2026-12-20
- GAP-009 (High): Document intended purpose, context of use, users, affected parties and explicitly out-of-scope or foreseeable misuse. — EI-04; owner: Product Owner; due 2026-12-20
- GAP-010 (High): Add the system to the obligations register and obtain a legal review of applicable obligations by jurisdiction. — AO-07; owner: Legal / Compliance; due 2026-12-20
- GAP-011 (High): Add AI-specific change triggers (model, prompt, data, tool, vendor) to change management and require reassessment when they fire. — GL-04, EI-05; owner: Engineering / Change Advisory Board; due 2026-12-20
- GAP-012 (High): Produce an AI-specific threat model covering injection, poisoning, abuse and tool misuse, and track mitigations to closure. — IS-01; owner: CISO / Security Architecture; due 2026-12-20
- GAP-013 (High): Restrict models and components to an approved list with integrity checks and dependency scanning. — IS-10, EI-05; owner: Security Architecture; due 2026-12-20
- GAP-014 (High): Define relevant groups and fairness metrics, test against thresholds and record mitigations. — IS-05; owner: Responsible AI Lead / Data Science; due 2026-12-20
- GAP-015 (High): Complete a harm analysis and safety evaluation, including vulnerable-user scenarios, and obtain residual-risk sign-off. — IS-06; owner: Responsible AI Lead; due 2026-12-20
- GAP-016 (High): Add clear AI-interaction and AI-generated-content disclosures at the point of interaction; have Legal review the wording. — IS-08; owner: Product Owner / Legal; due 2026-12-20
- GAP-017 (High): Deliver role-based training for operators and reviewers covering limitations, automation bias and escalation, and track completion. — AO-05; owner: HR / Learning lead with AI Governance Lead; due 2026-12-20

- GAP-018 (High): Define and communicate an escalation path and incident procedure for this system, with named contacts and timeframes. — AO-06, SE-04; owner: AI Governance Lead; due 2026-12-20
- GAP-019 (High): Define go-live conditions (controls, monitoring, oversight) and verify them with a signed checklist before production use. — GL-07; owner: Product Owner; due 2026-12-20
- GAP-020 (High): Negotiate or confirm AI contract clauses for data use, training restrictions, retention, breach notification and audit rights. — GL-11, AO-09; owner: Procurement / Vendor Risk; due 2026-12-20
- GAP-021 (High): Record the vendor AI product in the third-party AI inventory with an accountable owner and contract/DPA references. — EI-08, AO-09; owner: Procurement / Vendor Risk; due 2026-12-20
- GAP-022 (High): Monitor vendor model/version announcements and route changes through change control with reassessment triggers. — GL-04, EI-05; owner: Engineering / Change Advisory Board; due 2026-12-20
- GAP-023 (High): Bring the system under the AI incident procedure with defined severity levels and reporting timelines; Legal to confirm external reporting duties. — SE-04; owner: CISO / Operational Risk; due 2026-12-20
- GAP-024 (High): Store control evidence for this system in the evidence repository with retention rules and assemble an audit-ready evidence pack. — SE-02, SE-05; owner: AI Governance Lead; due 2026-12-20
- GAP-025 (High): Produce and maintain technical documentation including a system/model card, architecture, data, testing and known limitations. — GL-06; owner: Product / ML Engineering; due 2026-12-20
- GAP-026 (High): Schedule periodic review and recertification at the cadence set for the system's risk tier. — GL-12, EI-11; owner: AI Risk Lead; due 2026-12-20

Days 61–90 — standardise and evidence

- GAP-027 (Medium): Run the use case through intake triage and record a human-confirmed risk tier with rationale in the intake register. — GL-01, EI-02, EI-12; owner: AI Risk Lead; due 2027-01-19
- GAP-028 (Medium): Screen the use case against the prohibited/restricted AI use list and record the outcome; escalate any match for legal review. — AO-10; owner: Legal / AI Governance Lead; due 2027-01-19
- GAP-029 (Medium): Define and enforce retention periods for prompts, outputs, transcripts and logs consistent with the privacy assessment. — IS-03, SE-02; owner: DPO / Privacy Lead; due 2027-01-19
- GAP-030 (Medium): Register the agent and a tool/permission manifest listing each tool, scope and permitted action. — EI-07; owner: Platform Engineering Lead; due 2027-01-19
- GAP-031 (Medium): Replace shared/admin credentials with per-tool, least-privilege scoped identities and review access periodically. — IS-11; owner: Platform Engineering / IAM; due 2027-01-19
- GAP-032 (Medium): Document every data source, its classification and the data flows into prompts, retrieval indexes, logs and external services. — EI-06, IS-02; owner: Data Owner / CDO office; due 2027-01-19
- GAP-033 (Medium): Provide a route to a human, explanations proportionate to impact, and a contest/appeal channel whose outcomes are logged. — IS-08, AO-11; owner: Product Owner / Legal; due 2027-01-19

Beyond 90 days — evidence hygiene backlog

- GAP-034 (Low): Define acceptance criteria, run pre-release validation and robustness tests, and retain signed-off results. — GL-05, IS-07; owner: QA / ML Engineering; due 2027-04-19
- GAP-035 (Low): Enforce value caps, rate limits and step limits for agent actions in the runtime policy with alerts on breach. — IS-12, IS-16; owner: Platform Engineering / Security; due 2027-04-19
- GAP-036 (Low): Complete AI-specific vendor due diligence covering security, data use, training on customer data and subprocessors. — GL-11; owner: Procurement / Vendor Risk; due 2027-04-19

8. Audit-ready evidence checklist

Done	Question	Controls	Evidence expected	Evidence provided	Status
[ ]	UC-CORE-Q1	EI-03, EI-01	Inventory record with owner fields populated; Owner acceptance record	—	Missing evidence
[ ]	UC-CORE-Q2	EI-04	System profile with intended purpose; Users and affected parties; Foreseeable misuse analysis	—	Missing evidence
[ ]	UC-CORE-Q3	GL-01, EI-02, EI-12	Intake register entry; Triage decision with rationale; Tiering approval record	—	User-asserted
[ ]	UC-CORE-Q4	GL-02, GL-03	Completed risk/impact assessment; Reviewer sign-off; Approval record with conditions; Residual-risk statement	—	Missing evidence
[ ]	UC-CORE-Q5	AO-07	Obligations register entry for this system; Legal review record	—	Missing evidence
[ ]	UC-CORE-Q6	AO-10	Intake screening record against prohibited/restricted use list	—	User-asserted
[ ]	UC-CORE-Q7	GL-04, EI-05	Change policy with AI-specific triggers; Change tickets referencing reassessment; Model/prompt version history	—	Missing evidence
[ ]	UC-DATA-Q1	EI-06, IS-02	Data source register entries; Data-flow diagram; Data classification per source	—	User-asserted
[ ]	UC-DATA-Q2	IS-03	Privacy/data-protection impact assessment; Retention schedule; Transfer assessment	EVD-9749	Missing evidence
[ ]	UC-DATA-Q3	IS-04	Redaction configuration; Redaction test results; Log samples showing masking	—	Missing evidence
[ ]	UC-DATA-Q4	IS-03, IS-11	Access-scope configuration; Field-level access matrix; Access review record	EVD-8254	Missing evidence
[ ]	UC-DATA-Q5	IS-03, SE-02	Retention schedule covering AI artefacts; Deletion job configuration or records	—	User-asserted
[ ]	UC-SEC-Q1	IS-01	Threat model document; Mitigation tracking	—	Missing evidence
[ ]	UC-SEC-Q2	IS-10, EI-05	Approved model/provider list; Integrity checks; Dependency scanning results	—	Missing evidence
[x]	UC-SEC-Q3	IS-15	Kill-switch procedure; Kill-switch test record with time-to-disable; Rollback runbook	EVD-6101	Verified
[ ]	UC-SEC-Q4	GL-05, IS-07	Test plan and results; Acceptance criteria; Release sign-off	EVD-2824	User-asserted
[ ]	UC-AGT-Q1	EI-07	Agent registry entry; Tool/permission manifest per agent	—	User-asserted
[ ]	UC-AGT-Q2	IS-11	Tool permission matrix; Access-scope configuration; Access review records	—	User-asserted
[ ]	UC-AGT-Q3	IS-13	Approval rules in runtime policy; Approval records with approver identity	—	Critical gap
[ ]	UC-AGT-Q4	SE-08	Log configuration; Tamper-evidence mechanism; Retention settings; Sample log extract	EVD-1335	Missing evidence
[ ]	UC-AGT-Q5	IS-12, IS-16	Versioned runtime policy; Limit configuration; Limit-breach alerts	EVD-1618	User-asserted
[ ]	UC-FAIR-Q1	IS-05	Fairness test plan and results; Mitigation records	—	Missing evidence
[ ]	UC-FAIR-Q2	IS-06	Harm analysis; Safety evaluation results; Residual-risk sign-off	—	Missing evidence
[ ]	UC-FAIR-Q3	IS-08	Disclosure text and placement; User-facing documentation	—	Missing evidence
[ ]	UC-FAIR-Q4	IS-08, AO-11	Explanation approach; Feedback channel; Feedback log	—	User-asserted
[ ]	UC-HUM-Q1	IS-09	Oversight design document; Override/appeal procedures; Oversight operation records	—	Critical gap

Done	Question	Controls	Evidence expected	Evidence provided	Status
[ ]	UC-HUM-Q2	AO-05	Training curriculum; Completion records by role	—	Missing evidence
[ ]	UC-HUM-Q3	AO-06, SE-04	Escalation matrix; AI incident procedure; Sample escalation records	EVD-1642	Missing evidence
[ ]	UC-HUM-Q4	GL-07	Go-live checklist; Condition verification record	—	Missing evidence
[ ]	UC-VEN-Q1	GL-11	Vendor AI questionnaire; Due-diligence approval	EVD-3724	User-asserted
[ ]	UC-VEN-Q2	GL-11, AO-09	Contract AI clauses; Responsibility allocation (e.g. RACI or contract schedule)	—	Missing evidence
[ ]	UC-VEN-Q3	EI-08, AO-09	Vendor AI register; Vendor AI owner assignments	—	Missing evidence
[ ]	UC-VEN-Q4	GL-04, EI-05	Change tickets referencing reassessment; Model registry entries with versions	—	Missing evidence
[x]	UC-MON-Q1	GL-08, SE-12	Monitoring dashboards; Alert thresholds; Monitoring review records	EVD-2228	Verified
[ ]	UC-MON-Q2	SE-04	AI incident procedure; Incident log; Post-incident reviews	—	Missing evidence
[ ]	UC-MON-Q3	SE-02, SE-05	Evidence repository; Retention policy; Sample evidence pack	—	Missing evidence
[ ]	UC-MON-Q4	GL-12, EI-11	Review calendar by tier; Recertification records	—	Missing evidence
[ ]	UC-MON-Q5	GL-06	Technical documentation pack; Model/system card	—	Missing evidence

Only items marked [x] have evidence reviewed by a named reviewer.

9. Runtime guardrail recommendations

- **RG-PII-REDACTION** — Redact personal and sensitive data in prompts, responses and logs (IS-04, IS-03)
- **RG-TOOL-ALLOWLIST** — Allow-list tools and scopes per agent; deny by default (IS-11, EI-07)
- **RG-HUMAN-APPROVAL** — Require human approval for high-impact actions (refunds, record changes, external messages, deletions) (IS-13, IS-09)
- **RG-LIMITS** — Enforce rate, token, spend and per-action value limits (IS-16, IS-12)
- **RG-ACTION-LOGGING** — Log every model and tool call with trace ID, redacted inputs and policy decision (SE-08)
- **RG-KILL-SWITCH** — Provide a kill switch to disable the system, agent or individual tools (IS-15)
- **RG-DISCLOSURE** — Display AI-interaction disclosure to end users (IS-08)

AEGIS Shield advisor rating: High (see the draft runtime policy for details; activation in production requires human approval).

Draft guardrail policies can be generated with AEGIS Shield; activating or changing a production runtime policy requires human approval (IS-12, GL-04).

10. Human approvals

Approval	Type	Status	Resource	Requested by	Resolved by
apr_mvi8hs8jb290cdb0	confirm_risk_tier	approved	sys_mvi8hs7he2b521eb	u-assessor-1	u-approver-1

Appendix A — Evidence status definitions

Status	Definition
Verified	Evidence exists and has been reviewed
User-asserted	User claims implementation, but evidence is absent
Inferred	Inferred as a likely gap or fact from responses or metadata
Missing evidence	Required evidence has not been provided
Critical gap	Material privacy, security, safety, legal, or operational exposure

Appendix B — Audit trace

Tool calls and agent turns for this assessment are recorded in the hash-chained AEGIS audit log (SE-08).

Timestamp (UTC)	Trace ID	Action	Actor	Outcome	Audit record
2026-10-21T14:55:23	trc_mvi8hudy7303d33f	agent.tool.start_assessment	u-assessor-1	success	aud_mvi8hue0634b551d
2026-10-21T14:55:23	trc_mvi8hue182fe3079	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hue10409dde4
2026-10-21T14:55:23	trc_mvi8hue2413dc96d	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hue21f058238
2026-10-21T14:55:23	trc_mvi8hue3e5da1bf5	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hue5e9e5bf5a
2026-10-21T14:55:23	trc_mvi8hue6d65aa589	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hue6cd73d503
2026-10-21T14:55:23	trc_mvi8hue7c5f5814d	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hue9ae971e40
2026-10-21T14:55:23	trc_mvi8hue9dbf49e8e	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huea3d6c515d
2026-10-21T14:55:23	trc_mvi8huea432c0ac1	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hueb9c0398aa
2026-10-21T14:55:23	trc_mvi8huec1688d55c	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hueda563a8ba
2026-10-21T14:55:23	trc_mvi8hued5ec1e8dc	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huef431e0c93
2026-10-21T14:55:23	trc_mvi8hueg85b5b540	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huehafe1c574
2026-10-21T14:55:23	trc_mvi8huei515511e3	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huejb992d453
2026-10-21T14:55:23	trc_mvi8huekfc468204	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huel5eb7be19
2026-10-21T14:55:23	trc_mvi8huel4b0c9763	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huemeb3991d3
2026-10-21T14:55:23	trc_mvi8huen39b1cc73	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hueo4d6ced2d
2026-10-21T14:55:23	trc_mvi8huepb2d7cf0a	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huer2247694b
2026-10-21T14:55:23	trc_mvi8huer19fa060b	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huewd0388be5
2026-10-21T14:55:23	trc_mvi8huex99353b82	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huey9dbd5298
2026-10-21T14:55:23	trc_mvi8huey5f852d44	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huez4dbe5d1a
2026-10-21T14:55:23	trc_mvi8huf0de55b372	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huf0936c75f6
2026-10-21T14:55:23	trc_mvi8huf187ced4fa	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huf4361bb6f6

Timestamp (UTC)	Trace ID	Action	Actor	Outcome	Audit record
2026-10-21T14:55:23	trc_mvi8huf5cee16c57	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8huf8df3f9678
2026-10-21T14:55:23	trc_mvi8hufa3595373d	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufb2144454c
2026-10-21T14:55:23	trc_mvi8hufc66660efc	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufd12e07898
2026-10-21T14:55:23	trc_mvi8hufdf0553cdf	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufi630d056c
2026-10-21T14:55:23	trc_mvi8hufib82319da	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufk81ad442a
2026-10-21T14:55:23	trc_mvi8hufk742b1914	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufm1cfe96d6
2026-10-21T14:55:23	trc_mvi8hufm96c99483	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufoe04f23c8
2026-10-21T14:55:23	trc_mvi8hufp73688853	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufse27e658c
2026-10-21T14:55:23	trc_mvi8hufs3b0f5d72	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufuaf0f3d9f
2026-10-21T14:55:23	trc_mvi8hufu4b275ff7	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufx0a9dd333
2026-10-21T14:55:23	trc_mvi8hufxdaf689e2	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufyad98653f
2026-10-21T14:55:23	trc_mvi8hufy86851aaa	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hufze5f6057c
2026-10-21T14:55:23	trc_mvi8hug0093a63ca	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hug1ee390949
2026-10-21T14:55:23	trc_mvi8hug115f9e545	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hug335d3837c
2026-10-21T14:55:23	trc_mvi8hug43dba8253	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hug5850250a9
2026-10-21T14:55:23	trc_mvi8hug57599f80f	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hug6750d871e
2026-10-21T14:55:23	trc_mvi8hug7558e0272	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hugi71a79dde
2026-10-21T14:55:23	trc_mvi8hugi7e2b4fea	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8hugk3aa5499a
2026-10-21T14:55:23	trc_mvi8hugm670d6a11	evidence.reviewed	u-approver-1	success	aud_mvi8hugm2d2949ca
2026-10-21T14:55:23	trc_mvi8hugo24fe1ae0	evidence.reviewed	u-approver-1	success	aud_mvi8hugo8725b109
2026-10-21T14:55:23	trc_mvi8hugs3ef91b81	evidence.reviewed	u-approver-1	success	aud_mvi8hugs3b675b33
2026-10-21T14:55:23	trc_mvi8hugu1222c834	evidence.reviewed	u-approver-1	success	aud_mvi8huguaec6f054
2026-10-21T14:55:23	trc_mvi8hugwcbd17b83	evidence.reviewed	u-approver-1	success	aud_mvi8hugw5175454e
2026-10-21T14:55:23	trc_mvi8hugw8cc04288	agent.tool.identify_gaps	u-assessor-1	success	aud_mvi8huhf21a5dc10

Appendix C — Method and limitations

- Answers are self-reported; AEGIS has not inspected systems or documents unless evidence is marked Verified.

- The risk tier is preliminary and rule-based; the EU AI Act category screen is a trigger for legal review, not a legal classification.
- Framework mappings are indicative AEGIS interpretations from the Control Library, not official crosswalks.
- Regulatory applicability timelines must be verified against the current Official Journal text and any amendments.

Disclaimer

AEGIS supports governance readiness, control implementation and evidence collection. It does not certify ISO/IEC 42001 conformity, guarantee EU AI Act compliance, or replace legal advice. Framework mappings are indicative.

This use-case assessment produces an indicative view of governance readiness for a single AI system. Self-reported answers are User-asserted; only evidence reviewed by a named reviewer is Verified. Risk tiers are Inferred and require human confirmation. Framework mappings are indicative. AEGIS does not certify ISO/IEC 42001 conformity, guarantee EU AI Act compliance, or replace legal advice.

Synthetic example — fictional data.