AI Use-Case Governance Readiness Report — CV screening assistant

Synthetic example. *This report was generated from fictional data to illustrate AEGIS Agent output. It describes no real organization or system.*

Field	Value
Report ID	rpt_mvi8htun661d6a90
Assessment ID	asm_mvi8htshf3555a3a
AI system record	sys_mvi8hs7rbb9a6ca0
Generated (UTC)	2026-10-21T14:55:22+00:00
Prepared by	Synthetic AEGIS Analyst via AEGIS Agent
Question bank	AEGIS-UC-ASSESS v1.0.0
Control library	AEGIS Control Library v1.0.0
Evidence basis	Self-reported answers (User-asserted) unless evidence was reviewed by a named reviewer (Verified)

1. Executive summary

- **Preliminary risk tier: High** (Inferred; requires human confirmation — Confirmed by Synthetic Head of AI Risk (2026-10-21)).
- **38 of 38** applicable questions answered across 8 modules.
- **38 gaps** identified: 4 Critical, 23 High, 9 Medium, 2 Low.
- **Evidence posture:** 0 finding(s) Verified by a named reviewer; 10 User-asserted; 24 Missing evidence; 4 Critical gap; 0 Inferred.
- **Priority actions (Critical, due within 30 days):**
  - GAP-001: Register the system in the AI inventory with a named accountable business owner and technical owner, and capture owner acceptance. (EI-03, EI-01; owner: Business unit leaders)
  - GAP-002: Complete a privacy impact assessment (DPIA where required) covering lawful basis, minimisation, retention and transfers, signed off by the DPO/privacy lead. (IS-03; owner: DPO / Privacy Lead)
  - GAP-003: Define and enforce retention periods for prompts, outputs, transcripts and logs consistent with the privacy assessment. (IS-03, SE-02; owner: DPO / Privacy Lead)
  - GAP-004: Document the oversight design (reviewer roles, review points, override and stop authority) and keep records of oversight in operation. (IS-09; owner: Business Owner)

This report indicates governance readiness and evidence gaps for one AI system. It is not a determination of conformity or legal compliance with any framework or law.

2. AI system profile

Field	Value	Evidence status
name	CV screening assistant	User-asserted
purpose	Shortlist job applicants for recruiters	User-asserted
domain_category	employment	User-asserted
decision_impact	material	User-asserted
affected_parties	public	User-asserted
data_sensitivity	personal	User-asserted
autonomy	suggests	User-asserted
external_actions	No	User-asserted

Field	Value	Evidence status
genai	Yes	User-asserted
vendor_model	Yes	User-asserted
eu_exposure	Yes	User-asserted
tools	—	User-asserted
environment	production	User-asserted

3. Risk tier and rationale

Tier: High (score 100) — evidence status: Inferred; requires human confirmation: yes.

Rationale: Preliminary tier High (score 100). Drivers: decision_impact=material (+20); affected_parties=public (+15); data_sensitivity=personal (+15); autonomy=suggests (+5); domain_category=employment — Employment and worker management (Annex III area 4) (+30); generative_ai=True (+5); third_party_model=True (+5); eu_exposure=True (+5). Triggers: Potential EU AI Act high-risk area: Employment and worker management (Annex III area 4) — legal classification review required (Art. 6); Material or automated decisions in a sensitive (Annex III-like) domain. Inferred from user-provided intake; requires human confirmation.

Factor	Value	Points	Citation
decision_impact	material	20	AEGIS Control GL-02: Pre-deployment risk and impact assessment (Control Library v1.0.0)
affected_parties	public	15	AEGIS Control EI-04: Intended purpose and context of use (Control Library v1.0.0)
data_sensitivity	personal	15	AEGIS Control IS-03: Privacy assessment and data minimisation (Control Library v1.0.0)
autonomy	suggests	5	AEGIS Control IS-09: Human oversight design (Control Library v1.0.0)
domain_category	employment — Employment and worker management (Annex III area 4)	30	AEGIS Control AO-07: Regulatory and contractual obligations register (Control Library v1.0.0)
generative_ai	Yes	5	AEGIS Control IS-14: Prompt-injection and data-leakage testing (Control Library v1.0.0)
third_party_model	Yes	5	AEGIS Control GL-11: Vendor AI due diligence and contract controls (Control Library v1.0.0)
eu_exposure	Yes	5	AEGIS Control AO-07: Regulatory and contractual obligations register (Control Library v1.0.0)

Triggers:

- Potential EU AI Act high-risk area: Employment and worker management (Annex III area 4) — legal classification review required (Art. 6)
- Material or automated decisions in a sensitive (Annex III-like) domain
- Approval authority for this tier: AI Governance Forum
- Review cadence: every 6 months
- Assessment modules applied: Core governance & lifecycle, Data & privacy, Security & resilience, Generative AI & LLM controls, Fairness, safety & transparency, Human oversight & operations, Vendor & third-party AI, Monitoring, incidents & evidence
- Method: AEGIS risk-tier rules v1.0 (thresholds: High >= 60 or any high-risk trigger; Medium >= 30)

4. Domain heat map

Domain	Qs	Implementation (self-reported)	Verified	Crit	High	Med	Low	Rating
Strategy & Accountability	2	25.0%	0	0	2	0	0	High exposure

Domain	Qs	Implementation (self-reported)	Verified	Crit	High	Med	Low	Rating
AI Inventory & Visibility	3	0.0%	0	1	2	0	0	Critical exposure
Risk Management & Lifecycle	8	37.5%	0	0	6	2	0	High exposure
Data & Privacy	5	50.0%	0	2	1	2	0	Critical exposure
Security & Resilience	7	57.1%	0	0	4	2	1	High exposure
Fairness, Safety & Transparency	5	60.0%	0	0	2	2	1	High exposure
Human Oversight & Operations	5	30.0%	0	1	4	0	0	Critical exposure
Assurance, Evidence & Improvement	3	33.3%	0	0	2	1	0	High exposure

Implementation indicator = implemented (1.0) and partial (0.5) answers as a share of applicable questions; it reflects self-reported answers, not verified evidence.

5. Gap register

Gap · severity · due	Evidence status	Question	Controls	Owner role	Recommended action
GAP-001 · Critical · due 2026-11-20	Critical gap	UC-CORE-Q1: Is there a named accountable business owner and technical owner for this AI system, recorded in the AI inventory?	EI-03, EI-01	Business unit leaders	Register the system in the AI inventory with a named accountable business owner and technical owner, and capture owner acceptance.
GAP-002 · Critical · due 2026-11-20	Critical gap	UC-DATA-Q2: Has a privacy / data-protection impact assessment been completed covering lawful basis, minimisation, retention and cross-border transfers?	IS-03	DPO / Privacy Lead	Complete a privacy impact assessment (DPIA where required) covering lawful basis, minimisation, retention and transfers, signed off by the DPO/privacy lead.
GAP-003 · Critical · due 2026-11-20	Critical gap	UC-DATA-Q5: Are retention periods defined and enforced for prompts, outputs, conversation transcripts and logs?	IS-03, SE-02	DPO / Privacy Lead	Define and enforce retention periods for prompts, outputs, transcripts and logs consistent with the privacy assessment.
GAP-004 · Critical · due 2026-11-20	Critical gap	UC-HUM-Q1: Is the human oversight design documented — who reviews which outputs or actions, when, and with authority to override or stop the system?	IS-09	Business Owner	Document the oversight design (reviewer roles, review points, override and stop authority) and keep records of oversight in operation.
GAP-005 · High · due 2026-12-20	Missing evidence	UC-CORE-Q4: Was a pre-deployment risk and impact assessment completed and approved at the gate appropriate to the risk tier?	GL-02, GL-03	AI Risk Lead	Complete a tier-proportionate risk and impact assessment and obtain approval at the tier-appropriate gate before (further) production use.
GAP-006 · High · due 2026-12-20	Missing evidence	UC-CORE-Q6: Has the use case been screened against the organization's prohibited and restricted AI use policy?	AO-10	Legal / AI Governance Lead	Screen the use case against the prohibited/restricted AI use list and record the outcome; escalate any match for legal review.
GAP-007 · High · due 2026-12-20	Missing evidence	UC-GEN-Q1: Has the system been tested for prompt injection (direct and indirect), jailbreaks and data leakage before release and after material changes?	IS-14	Security Testing / Red Team	Run adversarial testing for direct and indirect prompt injection, jailbreaks and data leakage before release and after each material change.
GAP-008 · High · due 2026-12-20	Missing evidence	UC-CORE-Q2: Are the intended purpose, context of use, users, affected parties and out-of-scope uses documented?	EI-04	Product Owner	Document intended purpose, context of use, users, affected parties and explicitly out-of-scope or foreseeable misuse.

Gap · severity · due	Evidence status	Question	Controls	Owner role	Recommended action
GAP-009 · High · due 2026-12-20	Missing evidence	UC-CORE-Q5: Have the legal, regulatory and contractual obligations that apply to this system (by jurisdiction) been identified and reviewed?	AO-07	Legal / Compliance	Add the system to the obligations register and obtain a legal review of applicable obligations by jurisdiction.
GAP-010 · High · due 2026-12-20	Missing evidence	UC-CORE-Q7: Are changes to the model, prompts, data, tools or vendor subject to change control with defined reassessment triggers?	GL-04, EI-05	Engineering / Change Advisory Board	Add AI-specific change triggers (model, prompt, data, tool, vendor) to change management and require reassessment when they fire.
GAP-011 · High · due 2026-12-20	Missing evidence	UC-DATA-Q1: Are the data sources, classifications and data flows used for training, grounding (RAG) and prompts documented for this system?	EI-06, IS-02	Data Owner / CDO office	Document every data source, its classification and the data flows into prompts, retrieval indexes, logs and external services.
GAP-012 · High · due 2026-12-20	Missing evidence	UC-SEC-Q1: Has an AI-specific threat model been completed for this system (e.g. prompt injection, data poisoning, model abuse, tool misuse)?	IS-01	CISO / Security Architecture	Produce an AI-specific threat model covering injection, poisoning, abuse and tool misuse, and track mitigations to closure.
GAP-013 · High · due 2026-12-20	Missing evidence	UC-SEC-Q4: Were pre-release validation and robustness tests run against documented acceptance criteria, with results retained?	GL-05, IS-07	QA / ML Engineering	Define acceptance criteria, run pre-release validation and robustness tests, and retain signed-off results.
GAP-014 · High · due 2026-12-20	Missing evidence	UC-GEN-Q2: Are runtime guardrails (input/output filtering, PII detection, topic restrictions) enforced through a versioned runtime policy?	IS-12	Platform Engineering / Security	Enforce a versioned runtime policy for input/output filtering, PII detection and topic restrictions, starting in monitor mode then enforce.
GAP-015 · High · due 2026-12-20	Missing evidence	UC-GEN-Q4: Are system prompts and prompt templates version-controlled in a registry alongside the model version?	EI-05, IS-10	ML / Platform Engineering Lead	Version system prompts and templates in a registry linked to model versions and change tickets.
GAP-016 · High · due 2026-12-20	Missing evidence	UC-FAIR-Q2: Has a safety and harm evaluation been completed (harmful, misleading or inappropriate outputs, including for vulnerable users)?	IS-06	Responsible AI Lead	Complete a harm analysis and safety evaluation, including vulnerable-user scenarios, and obtain residual-risk sign-off.
GAP-017 · High · due 2026-12-20	Missing evidence	UC-FAIR-Q3: Are people informed when they are interacting with AI or receiving AI-generated content?	IS-08	Product Owner / Legal	Add clear AI-interaction and AI-generated-content disclosures at the point of interaction; have Legal review the wording.
GAP-018 · High · due 2026-12-20	Missing evidence	UC-HUM-Q2: Are operators and reviewers trained on the system's limitations, automation bias and escalation paths?	AO-05	HR / Learning lead with AI Governance Lead	Deliver role-based training for operators and reviewers covering limitations, automation bias and escalation, and track completion.
GAP-019 · High · due 2026-12-20	Missing evidence	UC-HUM-Q3: Is there a defined escalation path for concerns and AI incidents involving this system?	AO-06, SE-04	AI Governance Lead	Define and communicate an escalation path and incident procedure for this system, with named contacts and timeframes.
GAP-020 · High · due 2026-12-20	Missing evidence	UC-HUM-Q4: Were go-live deployment conditions defined, verified and signed off before production use?	GL-07	Product Owner	Define go-live conditions (controls, monitoring, oversight) and verify them with a signed checklist before production use.
GAP-021 · High · due 2026-12-20	Missing evidence	UC-VEN-Q2: Do contract terms cover data-use restrictions, no training on your data, retention, breach notification and audit rights?	GL-11, AO-09	Procurement / Vendor Risk	Negotiate or confirm AI contract clauses for data use, training restrictions, retention, breach notification and audit rights.

Gap · severity · due	Evidence status	Question	Controls	Owner role	Recommended action
GAP-022 · High · due 2026-12-20	Missing evidence	UC-VEN-Q3: Is the vendor AI product recorded in the third-party AI inventory with an accountable owner?	EI-08, AO-09	Procurement / Vendor Risk	Record the vendor AI product in the third-party AI inventory with an accountable owner and contract/DPA references.
GAP-023 · High · due 2026-12-20	Missing evidence	UC-VEN-Q4: Are vendor model version changes and deprecations monitored and routed through change control?	GL-04, EI-05	Engineering / Change Advisory Board	Monitor vendor model/version announcements and route changes through change control with reassessment triggers.
GAP-024 · High · due 2026-12-20	Missing evidence	UC-MON-Q1: Is production performance, drift and misuse monitored against thresholds with alerts and review?	GL-08, SE-12	ML Ops / Product Owner	Instrument performance, drift and misuse monitoring with thresholds, alerts and a recurring review.
GAP-025 · High · due 2026-12-20	Missing evidence	UC-MON-Q2: Is this system covered by an AI incident management procedure, including internal and external reporting timelines?	SE-04	CISO / Operational Risk	Bring the system under the AI incident procedure with defined severity levels and reporting timelines; Legal to confirm external reporting duties.
GAP-026 · High · due 2026-12-20	Missing evidence	UC-MON-Q3: Is evidence for this system's controls retained in an evidence repository, and can an audit-ready evidence pack be produced on request?	SE-02, SE-05	AI Governance Lead	Store control evidence for this system in the evidence repository with retention rules and assemble an audit-ready evidence pack.
GAP-027 · High · due 2026-12-20	Missing evidence	UC-MON-Q4: Is a periodic review / recertification of this system scheduled at the cadence required by its risk tier?	GL-12, EI-11	AI Risk Lead	Schedule periodic review and recertification at the cadence set for the system's risk tier.
GAP-028 · Medium · due 2027-01-19	User-asserted	UC-CORE-Q3: Has the use case passed intake triage and received a risk tier that was confirmed by an accountable approver?	GL-01, EI-02, EI-12	AI Risk Lead	Run the use case through intake triage and record a human-confirmed risk tier with rationale in the intake register.
GAP-029 · Medium · due 2027-01-19	User-asserted	UC-DATA-Q3: Are personal and sensitive fields redacted or masked before they reach prompts, logs and external model calls?	IS-04	Platform Engineering / Privacy Lead	Deploy redaction/masking on prompts, responses, logs and external model calls; test it against representative records and retain the results.
GAP-030 · Medium · due 2027-01-19	User-asserted	UC-DATA-Q4: Is the system's data access limited to the minimum records and fields it needs (e.g. field-level scoping of CRM access)?	IS-03, IS-11	DPO / Privacy Lead	Restrict data access to the minimum records and fields required, using scoped service identities, and review access periodically.
GAP-031 · Medium · due 2027-01-19	User-asserted	UC-SEC-Q2: Do models, prompts, plugins and third-party components come from an approved, integrity-checked supply chain?	IS-10, EI-05	Security Architecture	Restrict models and components to an approved list with integrity checks and dependency scanning.
GAP-032 · Medium · due 2027-01-19	User-asserted	UC-GEN-Q5: Are groundedness / hallucination rates and output quality evaluated against defined thresholds?	IS-07, IS-06	ML Engineering	Define groundedness and quality thresholds and evaluate against a representative test set before release and periodically.
GAP-033 · Medium · due 2027-01-19	User-asserted	UC-FAIR-Q4: Can affected individuals obtain an explanation, reach a human, and contest or appeal outcomes?	IS-08, AO-11	Product Owner / Legal	Provide a route to a human, explanations proportionate to impact, and a contest/appeal channel whose outcomes are logged.
GAP-034 · Medium · due 2027-01-19	User-asserted	UC-VEN-Q1: Was vendor AI due diligence completed for the model/service provider (security, data use, model training on your data, subprocessors)?	GL-11	Procurement / Vendor Risk	Complete AI-specific vendor due diligence covering security, data use, training on customer data and subprocessors.

Gap · severity · due	Evidence status	Question	Controls	Owner role	Recommended action
GAP-035 · Medium · due 2027-01-19	User-asserted	UC-MON-Q5: Is technical documentation (system/model card, architecture, data, testing, limitations) maintained for this system?	GL-06	Product / ML Engineering	Produce and maintain technical documentation including a system/model card, architecture, data, testing and known limitations.
GAP-036 · Medium · due 2027-01-19	Missing evidence	UC-GEN-Q3: Are rate, token and spend limits configured and alerting when breached?	IS-16	Platform Engineering / FinOps	Configure rate, token and budget limits per user and per system with breach alerts.
GAP-037 · Low · due 2027-04-19	User-asserted	UC-SEC-Q3: Is there a tested kill switch and rollback procedure that can disable the system, model, agent or individual tools quickly?	IS-15	Platform Engineering / SRE	Implement and test a kill switch and rollback runbook; record time-to-disable from the test.
GAP-038 · Low · due 2027-04-19	User-asserted	UC-FAIR-Q1: Has the system been tested for bias or discriminatory outcomes across relevant groups, with documented thresholds?	IS-05	Responsible AI Lead / Data Science	Define relevant groups and fairness metrics, test against thresholds and record mitigations.

Due dates apply the AEGIS severity SLAs (Critical 30, High 60, Medium 90, Low 180 days) from 2026-10-21.

6. Control recommendations

Control	Title	Priority	Highest gap severity	Owner role	Framework mappings (indicative)
EI-01	Enterprise AI inventory (system of record)	critical	Critical	AI Governance Lead / CIO office	NIST AI RMF: GOVERN 1.6; ISO/IEC 42001: A.4.2, 7.5; EU AI Act: Art. 26, Art. 49
EI-03	Accountable owner per AI system	critical	Critical	Business unit leaders	NIST AI RMF: GOVERN 2.1; ISO/IEC 42001: A.3.2, A.10.2; EU AI Act: Art. 26
IS-03	Privacy assessment and data minimisation	critical	Critical	DPO / Privacy Lead	NIST AI RMF: MEASURE 2.10, MAP 1.1; ISO/IEC 42001: A.7.4, A.5.4; EU AI Act: Art. 10
IS-09	Human oversight design	critical	Critical	Business Owner	NIST AI RMF: GOVERN 3.2, MAP 3.5; ISO/IEC 42001: A.9.2, A.6.2.6; EU AI Act: Art. 14, Art. 26
SE-02	Evidence repository and retention	high	Critical	AI Governance Lead	NIST AI RMF: GOVERN 1.4; ISO/IEC 42001: 7.5; EU AI Act: Art. 18, Art. 19
AO-10	Prohibited and restricted AI use policy	critical	High	Legal / AI Governance Lead	NIST AI RMF: GOVERN 1.1, MAP 1.1; ISO/IEC 42001: A.9.4, A.2.2; EU AI Act: Art. 5
GL-02	Pre-deployment risk and impact assessment	critical	High	AI Risk Lead	NIST AI RMF: MAP 1.1, MAP 3.1, MAP 5.1, MEASURE 1.1; ISO/IEC 42001: 6.1.4, 8.4, A.5.2, A.5.3, A.5.4; EU AI Act: Art. 9, Art. 27
GL-03	Risk-tiered approval gates	critical	High	AI Governance Forum	NIST AI RMF: MANAGE 1.1, GOVERN 1.4; ISO/IEC 42001: 8.1, A.6.2.5; EU AI Act: Art. 9, Art. 43
IS-14	Prompt-injection and data-leakage testing	critical	High	Security Testing / Red Team	NIST AI RMF: MEASURE 2.7; ISO/IEC 42001: A.6.2.4; EU AI Act: Art. 15
AO-05	AI literacy and role-based training	high	High	HR / Learning lead with AI Governance Lead	NIST AI RMF: GOVERN 2.2; ISO/IEC 42001: 7.2, 7.3, A.4.6; EU AI Act: Art. 4
AO-06	Escalation paths and decision rights	high	High	AI Governance Lead	NIST AI RMF: GOVERN 2.1, GOVERN 4.3; ISO/IEC 42001: 5.3, A.3.3; EU AI Act: Art. 26
AO-07	Regulatory and contractual obligations register	high	High	Legal / Compliance	NIST AI RMF: GOVERN 1.1; ISO/IEC 42001: 4.1, 4.2, 6.1; EU AI Act: Art. 6, Art. 26

Control	Title	Priority	Highest gap severity	Owner role	Framework mappings (indicative)
EI-04	Intended purpose and context of use	high	High	Product Owner	NIST AI RMF: MAP 1.1, MAP 1.2, MAP 3.4; ISO/IEC 42001: A.6.2.2, A.9.4; EU AI Act: Art. 11, Art. 13
EI-05	Model, prompt and component registry	high	High	ML / Platform Engineering Lead	NIST AI RMF: MAP 4.1, GOVERN 1.6; ISO/IEC 42001: A.4.4, A.6.2.3; EU AI Act: Art. 11
EI-06	Data source and lineage registry	high	High	Data Owner / CDO office	NIST AI RMF: MAP 2.3, MEASURE 2.10; ISO/IEC 42001: A.4.3, A.7.5; EU AI Act: Art. 10
EI-08	Third-party and vendor AI inventory	high	High	Procurement / Vendor Risk	NIST AI RMF: GOVERN 6.1; ISO/IEC 42001: A.10.3; EU AI Act: Art. 25
GL-04	Change management and reassessment triggers	high	High	Engineering / Change Advisory Board	NIST AI RMF: GOVERN 1.4, MANAGE 2.4; ISO/IEC 42001: 6.3, A.6.2.6; EU AI Act: Art. 43
GL-05	Pre-release validation and testing	high	High	QA / ML Engineering	NIST AI RMF: MEASURE 2.1, MEASURE 2.3, MEASURE 2.5; ISO/IEC 42001: A.6.2.4; EU AI Act: Art. 9, Art. 15
GL-07	Deployment conditions and go-live readiness	high	High	Product Owner	NIST AI RMF: MANAGE 1.2, MANAGE 2.1; ISO/IEC 42001: A.6.2.5; EU AI Act: Art. 26
GL-08	Post-deployment monitoring	high	High	ML Ops / Product Owner	NIST AI RMF: MEASURE 3.1, MANAGE 4.1; ISO/IEC 42001: 9.1, A.6.2.6; EU AI Act: Art. 26, Art. 72
GL-11	Vendor AI due diligence and contract controls	high	High	Procurement / Vendor Risk	NIST AI RMF: GOVERN 6.1, MAP 4.1; ISO/IEC 42001: A.10.3, A.10.4; EU AI Act: Art. 25
IS-01	AI threat modelling	high	High	CISO / Security Architecture	NIST AI RMF: MAP 5.1, MEASURE 2.7; ISO/IEC 42001: 6.1.2, A.6.1.3; EU AI Act: Art. 15
IS-02	Data governance and quality	high	High	Data Owner	NIST AI RMF: MAP 2.3, MEASURE 2.1; ISO/IEC 42001: A.7.2, A.7.3, A.7.4, A.7.6; EU AI Act: Art. 10
IS-06	Safety and harm evaluation	high	High	Responsible AI Lead	NIST AI RMF: MEASURE 2.6, MAP 5.1; ISO/IEC 42001: A.5.4, A.5.5; EU AI Act: Art. 9
IS-07	Accuracy and robustness testing	high	High	ML Engineering	NIST AI RMF: MEASURE 2.5; ISO/IEC 42001: A.6.2.4; EU AI Act: Art. 15
IS-08	Transparency and AI disclosures	high	High	Product Owner / Legal	NIST AI RMF: MEASURE 2.8, MEASURE 2.9; ISO/IEC 42001: A.8.2, A.8.5; EU AI Act: Art. 13, Art. 50
IS-10	Secure model and prompt supply chain	high	High	Security Architecture	NIST AI RMF: GOVERN 6.1, MEASURE 2.7; ISO/IEC 42001: A.10.3, A.6.2.3; EU AI Act: Art. 15
IS-12	Runtime guardrails and policy enforcement	high	High	Platform Engineering / Security	NIST AI RMF: MANAGE 2.2, MANAGE 2.4, MEASURE 2.7; ISO/IEC 42001: A.6.2.6; EU AI Act: Art. 15, Art. 14
SE-04	AI incident management and reporting	high	High	CISO / Operational Risk	NIST AI RMF: MANAGE 4.3, MANAGE 2.3; ISO/IEC 42001: 10.2, A.8.4; EU AI Act: Art. 73, Art. 26
SE-05	Audit-ready evidence packs	high	High	AI Governance Lead	NIST AI RMF: GOVERN 1.4, MEASURE 2.9; ISO/IEC 42001: 7.5, 9.2; EU AI Act: Art. 11, Art. 18
SE-12	Model performance and drift monitoring	high	High	ML Ops	NIST AI RMF: MEASURE 2.4, MEASURE 3.1; ISO/IEC 42001: A.6.2.6; EU AI Act: Art. 72, Art. 15
AO-09	Third-party AI accountability	medium	High	Procurement / Vendor Risk	NIST AI RMF: GOVERN 6.1, GOVERN 6.2; ISO/IEC 42001: A.10.2, A.10.3; EU AI Act: Art. 25, Art. 26

Control	Title	Priority	Highest gap severity	Owner role	Framework mappings (indicative)
EI-11	Inventory attestation and recertification	medium	High	AI Governance Lead	NIST AI RMF: GOVERN 1.5; ISO/IEC 42001: 9.1
GL-12	Periodic review and recertification	medium	High	AI Risk Lead	NIST AI RMF: GOVERN 1.5; ISO/IEC 42001: 9.3; EU AI Act: Art. 9
EI-02	AI risk classification and tiering	critical	Medium	AI Risk Lead	NIST AI RMF: MAP 1.1, MAP 1.5, GOVERN 1.3; ISO/IEC 42001: 6.1.2, A.5.2; EU AI Act: Art. 6
IS-04	Sensitive data redaction and masking	critical	Medium	Platform Engineering / Privacy Lead	NIST AI RMF: MEASURE 2.10; ISO/IEC 42001: A.7.4, A.6.2.6; EU AI Act: Art. 10
IS-11	Least-privilege tool access	critical	Medium	Platform Engineering / IAM	NIST AI RMF: MEASURE 2.7, MANAGE 2.2; ISO/IEC 42001: A.6.2.6, A.4.4; EU AI Act: Art. 15, Art. 14
EI-12	Use-case intake register	high	Medium	AI CoE	NIST AI RMF: MAP 1.1; ISO/IEC 42001: A.5.2, A.6.1.3; EU AI Act: Art. 6
GL-01	Intake triage workflow	high	Medium	AI CoE	NIST AI RMF: MAP 1.1, MAP 1.6; ISO/IEC 42001: 8.1, A.6.1.3; EU AI Act: Art. 9
GL-06	Technical documentation	high	Medium	Product / ML Engineering	NIST AI RMF: MAP 2.2, MEASURE 2.9; ISO/IEC 42001: 7.5, A.6.2.7, A.6.2.3; EU AI Act: Art. 11
AO-11	Stakeholder engagement and feedback	medium	Medium	Product Owner	NIST AI RMF: GOVERN 5.1, GOVERN 5.2; ISO/IEC 42001: 4.2, A.8.3, A.8.5; EU AI Act: Art. 27
IS-16	Rate, budget and resource limits	medium	Medium	Platform Engineering / FinOps	NIST AI RMF: MANAGE 2.2; ISO/IEC 42001: A.4.5, A.6.2.6; EU AI Act: Art. 15
IS-15	Kill switch and rollback	critical	Low	Platform Engineering / SRE	NIST AI RMF: MANAGE 2.4; ISO/IEC 42001: A.6.2.6; EU AI Act: Art. 14
IS-05	Bias and fairness testing	high	Low	Responsible AI Lead / Data Science	NIST AI RMF: MEASURE 2.11; ISO/IEC 42001: A.5.4, A.7.4; EU AI Act: Art. 10

7. Remediation roadmap (30/60/90 days)

First 30 days — contain Critical exposure

- GAP-001 (Critical): Register the system in the AI inventory with a named accountable business owner and technical owner, and capture owner acceptance. — EI-03, EI-01; owner: Business unit leaders; due 2026-11-20
- GAP-002 (Critical): Complete a privacy impact assessment (DPIA where required) covering lawful basis, minimisation, retention and transfers, signed off by the DPO/privacy lead. — IS-03; owner: DPO / Privacy Lead; due 2026-11-20
- GAP-003 (Critical): Define and enforce retention periods for prompts, outputs, transcripts and logs consistent with the privacy assessment. — IS-03, SE-02; owner: DPO / Privacy Lead; due 2026-11-20
- GAP-004 (Critical): Document the oversight design (reviewer roles, review points, override and stop authority) and keep records of oversight in operation. — IS-09; owner: Business Owner; due 2026-11-20

Days 31–60 — close High gaps

- GAP-005 (High): Complete a tier-proportionate risk and impact assessment and obtain approval at the tier-appropriate gate before (further) production use. — GL-02, GL-03; owner: AI Risk Lead; due 2026-12-20
- GAP-006 (High): Screen the use case against the prohibited/restricted AI use list and record the outcome; escalate any match for legal review. — AO-10; owner: Legal / AI Governance Lead; due 2026-12-20
- GAP-007 (High): Run adversarial testing for direct and indirect prompt injection, jailbreaks and data leakage before release and after each material change. — IS-14; owner: Security Testing / Red Team; due 2026-12-20
- GAP-008 (High): Document intended purpose, context of use, users, affected parties and explicitly out-of-scope or foreseeable misuse. — EI-04; owner: Product Owner; due 2026-12-20

- GAP-009 (High): Add the system to the obligations register and obtain a legal review of applicable obligations by jurisdiction. — AO-07; owner: Legal / Compliance; due 2026-12-20
- GAP-010 (High): Add AI-specific change triggers (model, prompt, data, tool, vendor) to change management and require reassessment when they fire. — GL-04, EI-05; owner: Engineering / Change Advisory Board; due 2026-12-20
- GAP-011 (High): Document every data source, its classification and the data flows into prompts, retrieval indexes, logs and external services. — EI-06, IS-02; owner: Data Owner / CDO office; due 2026-12-20
- GAP-012 (High): Produce an AI-specific threat model covering injection, poisoning, abuse and tool misuse, and track mitigations to closure. — IS-01; owner: CISO / Security Architecture; due 2026-12-20
- GAP-013 (High): Define acceptance criteria, run pre-release validation and robustness tests, and retain signed-off results. — GL-05, IS-07; owner: QA / ML Engineering; due 2026-12-20
- GAP-014 (High): Enforce a versioned runtime policy for input/output filtering, PII detection and topic restrictions, starting in monitor mode then enforce. — IS-12; owner: Platform Engineering / Security; due 2026-12-20
- GAP-015 (High): Version system prompts and templates in a registry linked to model versions and change tickets. — EI-05, IS-10; owner: ML / Platform Engineering Lead; due 2026-12-20
- GAP-016 (High): Complete a harm analysis and safety evaluation, including vulnerable-user scenarios, and obtain residual-risk sign-off. — IS-06; owner: Responsible AI Lead; due 2026-12-20
- GAP-017 (High): Add clear AI-interaction and AI-generated-content disclosures at the point of interaction; have Legal review the wording. — IS-08; owner: Product Owner / Legal; due 2026-12-20
- GAP-018 (High): Deliver role-based training for operators and reviewers covering limitations, automation bias and escalation, and track completion. — AO-05; owner: HR / Learning lead with AI Governance Lead; due 2026-12-20
- GAP-019 (High): Define and communicate an escalation path and incident procedure for this system, with named contacts and timeframes. — AO-06, SE-04; owner: AI Governance Lead; due 2026-12-20
- GAP-020 (High): Define go-live conditions (controls, monitoring, oversight) and verify them with a signed checklist before production use. — GL-07; owner: Product Owner; due 2026-12-20
- GAP-021 (High): Negotiate or confirm AI contract clauses for data use, training restrictions, retention, breach notification and audit rights. — GL-11, AO-09; owner: Procurement / Vendor Risk; due 2026-12-20
- GAP-022 (High): Record the vendor AI product in the third-party AI inventory with an accountable owner and contract/DPA references. — EI-08, AO-09; owner: Procurement / Vendor Risk; due 2026-12-20
- GAP-023 (High): Monitor vendor model/version announcements and route changes through change control with reassessment triggers. — GL-04, EI-05; owner: Engineering / Change Advisory Board; due 2026-12-20
- GAP-024 (High): Instrument performance, drift and misuse monitoring with thresholds, alerts and a recurring review. — GL-08, SE-12; owner: ML Ops / Product Owner; due 2026-12-20
- GAP-025 (High): Bring the system under the AI incident procedure with defined severity levels and reporting timelines; Legal to confirm external reporting duties. — SE-04; owner: CISO / Operational Risk; due 2026-12-20
- GAP-026 (High): Store control evidence for this system in the evidence repository with retention rules and assemble an audit-ready evidence pack. — SE-02, SE-05; owner: AI Governance Lead; due 2026-12-20
- GAP-027 (High): Schedule periodic review and recertification at the cadence set for the system's risk tier. — GL-12, EI-11; owner: AI Risk Lead; due 2026-12-20

Days 61–90 — standardise and evidence

- GAP-028 (Medium): Run the use case through intake triage and record a human-confirmed risk tier with rationale in the intake register. — GL-01, EI-02, EI-12; owner: AI Risk Lead; due 2027-01-19
- GAP-029 (Medium): Deploy redaction/masking on prompts, responses, logs and external model calls; test it against representative records and retain the results. — IS-04; owner: Platform Engineering / Privacy Lead; due 2027-01-19

- GAP-030 (Medium): Restrict data access to the minimum records and fields required, using scoped service identities, and review access periodically. — IS-03, IS-11; owner: DPO / Privacy Lead; due 2027-01-19
- GAP-031 (Medium): Restrict models and components to an approved list with integrity checks and dependency scanning. — IS-10, EI-05; owner: Security Architecture; due 2027-01-19
- GAP-032 (Medium): Define groundedness and quality thresholds and evaluate against a representative test set before release and periodically. — IS-07, IS-06; owner: ML Engineering; due 2027-01-19
- GAP-033 (Medium): Provide a route to a human, explanations proportionate to impact, and a contest/appeal channel whose outcomes are logged. — IS-08, AO-11; owner: Product Owner / Legal; due 2027-01-19
- GAP-034 (Medium): Complete AI-specific vendor due diligence covering security, data use, training on customer data and subprocessors. — GL-11; owner: Procurement / Vendor Risk; due 2027-01-19
- GAP-035 (Medium): Produce and maintain technical documentation including a system/model card, architecture, data, testing and known limitations. — GL-06; owner: Product / ML Engineering; due 2027-01-19
- GAP-036 (Medium): Configure rate, token and budget limits per user and per system with breach alerts. — IS-16; owner: Platform Engineering / FinOps; due 2027-01-19

Beyond 90 days — evidence hygiene backlog

- GAP-037 (Low): Implement and test a kill switch and rollback runbook; record time-to-disable from the test. — IS-15; owner: Platform Engineering / SRE; due 2027-04-19
- GAP-038 (Low): Define relevant groups and fairness metrics, test against thresholds and record mitigations. — IS-05; owner: Responsible AI Lead / Data Science; due 2027-04-19

8. Audit-ready evidence checklist

Done	Question	Controls	Evidence expected	Evidence provided	Status
[ ]	UC-CORE-Q1	EI-03, EI-01	Inventory record with owner fields populated; Owner acceptance record	—	Critical gap
[ ]	UC-CORE-Q2	EI-04	System profile with intended purpose; Users and affected parties; Foreseeable misuse analysis	—	Missing evidence
[ ]	UC-CORE-Q3	GL-01, EI-02, EI-12	Intake register entry; Triage decision with rationale; Tiering approval record	—	User-asserted
[ ]	UC-CORE-Q4	GL-02, GL-03	Completed risk/impact assessment; Reviewer sign-off; Approval record with conditions; Residual-risk statement	—	Missing evidence
[ ]	UC-CORE-Q5	AO-07	Obligations register entry for this system; Legal review record	—	Missing evidence
[ ]	UC-CORE-Q6	AO-10	Intake screening record against prohibited/restricted use list	EVD-1054	Missing evidence
[ ]	UC-CORE-Q7	GL-04, EI-05	Change policy with AI-specific triggers; Change tickets referencing reassessment; Model/prompt version history	—	Missing evidence
[ ]	UC-DATA-Q1	EI-06, IS-02	Data source register entries; Data-flow diagram; Data classification per source	—	Missing evidence
[ ]	UC-DATA-Q2	IS-03	Privacy/data-protection impact assessment; Retention schedule; Transfer assessment	—	Critical gap
[ ]	UC-DATA-Q3	IS-04	Redaction configuration; Redaction test results; Log samples showing masking	—	User-asserted
[ ]	UC-DATA-Q4	IS-03, IS-11	Access-scope configuration; Field-level access matrix; Access review record	—	User-asserted
[ ]	UC-DATA-Q5	IS-03, SE-02	Retention schedule covering AI artefacts; Deletion job configuration or records	—	Critical gap
[ ]	UC-SEC-Q1	IS-01	Threat model document; Mitigation tracking	—	Missing evidence
[ ]	UC-SEC-Q2	IS-10, EI-05	Approved model/provider list; Integrity checks; Dependency scanning results	—	User-asserted

Done	Question	Controls	Evidence expected	Evidence provided	Status
[ ]	UC-SEC-Q3	IS-15	Kill-switch procedure; Kill-switch test record with time-to-disable; Rollback runbook	EVD-7014	User-asserted
[ ]	UC-SEC-Q4	GL-05, IS-07	Test plan and results; Acceptance criteria; Release sign-off	EVD-9445	Missing evidence
[ ]	UC-GEN-Q1	IS-14	Red-team / adversarial test results; Remediation records; Re-test evidence	—	Missing evidence
[ ]	UC-GEN-Q2	IS-12	Versioned runtime policy; Policy decision logs; Enforcement-mode configuration	—	Missing evidence
[ ]	UC-GEN-Q3	IS-16	Limit configuration; Limit-breach alerts	—	Missing evidence
[ ]	UC-GEN-Q4	EI-05, IS-10	Model registry entries with versions; Prompt/config version history	—	Missing evidence
[ ]	UC-GEN-Q5	IS-07, IS-06	Accuracy metrics with thresholds; Safety evaluation results	—	User-asserted
[ ]	UC-FAIR-Q1	IS-05	Fairness test plan and results; Mitigation records	EVD-4902	User-asserted
[ ]	UC-FAIR-Q2	IS-06	Harm analysis; Safety evaluation results; Residual-risk sign-off	—	Missing evidence
[ ]	UC-FAIR-Q3	IS-08	Disclosure text and placement; User-facing documentation	—	Missing evidence
[ ]	UC-FAIR-Q4	IS-08, AO-11	Explanation approach; Feedback channel; Feedback log	—	User-asserted
[ ]	UC-HUM-Q1	IS-09	Oversight design document; Override/appeal procedures; Oversight operation records	—	Critical gap
[ ]	UC-HUM-Q2	AO-05	Training curriculum; Completion records by role	—	Missing evidence
[ ]	UC-HUM-Q3	AO-06, SE-04	Escalation matrix; AI incident procedure; Sample escalation records	—	Missing evidence
[ ]	UC-HUM-Q4	GL-07	Go-live checklist; Condition verification record	EVD-4052	Missing evidence
[ ]	UC-VEN-Q1	GL-11	Vendor AI questionnaire; Due-diligence approval	—	User-asserted
[ ]	UC-VEN-Q2	GL-11, AO-09	Contract AI clauses; Responsibility allocation (e.g. RACI or contract schedule)	—	Missing evidence
[ ]	UC-VEN-Q3	EI-08, AO-09	Vendor AI register; Vendor AI owner assignments	—	Missing evidence
[ ]	UC-VEN-Q4	GL-04, EI-05	Change tickets referencing reassessment; Model registry entries with versions	—	Missing evidence
[ ]	UC-MON-Q1	GL-08, SE-12	Monitoring dashboards; Alert thresholds; Monitoring review records	—	Missing evidence
[ ]	UC-MON-Q2	SE-04	AI incident procedure; Incident log; Post-incident reviews	—	Missing evidence
[ ]	UC-MON-Q3	SE-02, SE-05	Evidence repository; Retention policy; Sample evidence pack	—	Missing evidence
[ ]	UC-MON-Q4	GL-12, EI-11	Review calendar by tier; Recertification records	—	Missing evidence
[ ]	UC-MON-Q5	GL-06	Technical documentation pack; Model/system card	—	User-asserted

Only items marked [x] have evidence reviewed by a named reviewer.

9. Runtime guardrail recommendations

- **RG-PII-REDACTION** — Redact personal and sensitive data in prompts, responses and logs (IS-04, IS-03)

- **RG-INJECTION-FILTER** — Detect and block prompt-injection and jailbreak attempts, including indirect injection in retrieved content (IS-14, IS-12)
- **RG-LIMITS** — Enforce rate, token, spend and per-action value limits (IS-16, IS-12)
- **RG-ACTION-LOGGING** — Log every model and tool call with trace ID, redacted inputs and policy decision (SE-08)
- **RG-KILL-SWITCH** — Provide a kill switch to disable the system, agent or individual tools (IS-15)
- **RG-DISCLOSURE** — Display AI-interaction disclosure to end users (IS-08)

AEGIS Shield advisor rating: High (see the draft runtime policy for details; activation in production requires human approval).

Draft guardrail policies can be generated with AEGIS Shield; activating or changing a production runtime policy requires human approval (IS-12, GL-04).

10. Human approvals

Approval	Type	Status	Resource	Requested by	Resolved by
apr_mvi8hs8z2fb0cd9b	confirm_risk_tier	approved	sys_mvi8hs7rbb9a6ca0	u-assessor-1	u-approver-1

Appendix A — Evidence status definitions

Status	Definition
Verified	Evidence exists and has been reviewed
User-asserted	User claims implementation, but evidence is absent
Inferred	Inferred as a likely gap or fact from responses or metadata
Missing evidence	Required evidence has not been provided
Critical gap	Material privacy, security, safety, legal, or operational exposure

Appendix B — Audit trace

Tool calls and agent turns for this assessment are recorded in the hash-chained AEGIS audit log (SE-08).

Timestamp (UTC)	Trace ID	Action	Actor	Outcome	Audit record
2026-10-21T14:55:22	trc_mvi8htsgdccb0576	agent.tool.start_assessment	u-assessor-1	success	aud_mvi8htsi096105ae
2026-10-21T14:55:22	trc_mvi8htsi2cb537cc	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsj219a4151
2026-10-21T14:55:22	trc_mvi8htsj0e967d88	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsk76034cbb
2026-10-21T14:55:22	trc_mvi8htsl791c4772	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsl835cde4a
2026-10-21T14:55:22	trc_mvi8htsm7879cf7d	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsm1b823bf6
2026-10-21T14:55:22	trc_mvi8htsn8ade89c0	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsnc390cf89
2026-10-21T14:55:22	trc_mvi8htsoad454b8f	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htspf7b069ef
2026-10-21T14:55:22	trc_mvi8htsqf57afc30	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsq41dcca55
2026-10-21T14:55:22	trc_mvi8htsrcff325a1	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htssd48a55c7
2026-10-21T14:55:22	trc_mvi8htssb7f5d799	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htstd46b8105

Timestamp (UTC)	Trace ID	Action	Actor	Outcome	Audit record
2026-10-21T14:55:22	trc_mvi8htste46cae45	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsub81c510d
2026-10-21T14:55:22	trc_mvi8htsua138b4d5	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsv589e508b
2026-10-21T14:55:22	trc_mvi8htsv4ff25b8e	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsw8fe53deb
2026-10-21T14:55:22	trc_mvi8htsxb7b5b914	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsyf75b6915
2026-10-21T14:55:22	trc_mvi8htsy60a4145d	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htsze676fab5
2026-10-21T14:55:22	trc_mvi8htszb78c073b	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htt11536f6d0
2026-10-21T14:55:22	trc_mvi8htt2781b8194	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htt42a7fdc86
2026-10-21T14:55:22	trc_mvi8htt49d39ce4a	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htt78139e52f
2026-10-21T14:55:22	trc_mvi8htt88b62920e	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httb81bd32c2
2026-10-21T14:55:22	trc_mvi8httb8fb2d802	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httd25421c30
2026-10-21T14:55:22	trc_mvi8httd2e9d2760	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htte5a2991b5
2026-10-21T14:55:22	trc_mvi8httff3481627	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httg0299010d
2026-10-21T14:55:22	trc_mvi8httgc306b5ee	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htti5ee28906
2026-10-21T14:55:22	trc_mvi8htti9e2d36c8	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httjbbeb7e6d
2026-10-21T14:55:22	trc_mvi8httje5d124af	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httk9a17107d
2026-10-21T14:55:22	trc_mvi8httl3690e974	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httme565d9a0
2026-10-21T14:55:22	trc_mvi8httmf1c00676	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8http4f38e177
2026-10-21T14:55:22	trc_mvi8httqc7c5e1ad	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httr5a16d0fb
2026-10-21T14:55:22	trc_mvi8httr298e4c33	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htts4617b88e
2026-10-21T14:55:22	trc_mvi8htts9bc01c3e	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httu8e48a968
2026-10-21T14:55:22	trc_mvi8httub5ba82c4	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httv7076ccd3
2026-10-21T14:55:22	trc_mvi8httw1189fd6c	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httxf675a8a2
2026-10-21T14:55:22	trc_mvi8httxf5cb455a	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8httya0a47004
2026-10-21T14:55:22	trc_mvi8httz3f722723	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htu06040e0a1
2026-10-21T14:55:22	trc_mvi8htu06f48cbc3	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htu13b006d6e
2026-10-21T14:55:22	trc_mvi8htu172f53374	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htu2d9c7bf35

Timestamp (UTC)	Trace ID	Action	Actor	Outcome	Audit record
2026-10-21T14:55:22	trc_mvi8htu3d04d444	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htu438d811ab
2026-10-21T14:55:22	trc_mvi8htu4ccc829ef	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htu539eed76d
2026-10-21T14:55:22	trc_mvi8htu5d11b13ff	agent.tool.submit_assessment_response	u-assessor-1	success	aud_mvi8htu6c387509
2026-10-21T14:55:22	trc_mvi8htu8c4c15c2e	evidence.reviewed	u-approver-1	success	aud_mvi8htu8f14f5e8f
2026-10-21T14:55:22	trc_mvi8htu92b60e3c0	evidence.reviewed	u-approver-1	success	aud_mvi8htu9598ab7d7
2026-10-21T14:55:22	trc_mvi8htub3f8e8868	evidence.reviewed	u-approver-1	success	aud_mvi8htub62e78aa4
2026-10-21T14:55:22	trc_mvi8htub6b01b4e6	agent.tool.identify_gaps	u-assessor-1	success	aud_mvi8htuf8e062cb8

Appendix C — Method and limitations

- Answers are self-reported; AEGIS has not inspected systems or documents unless evidence is marked Verified.
- The risk tier is preliminary and rule-based; the EU AI Act category screen is a trigger for legal review, not a legal classification.
- Framework mappings are indicative AEGIS interpretations from the Control Library, not official crosswalks.
- Regulatory applicability timelines must be verified against the current Official Journal text and any amendments.

Disclaimer

AEGIS supports governance readiness, control implementation and evidence collection. It does not certify ISO/IEC 42001 conformity, guarantee EU AI Act compliance, or replace legal advice. Framework mappings are indicative.

This use-case assessment produces an indicative view of governance readiness for a single AI system. Self-reported answers are User-asserted; only evidence reviewed by a named reviewer is Verified. Risk tiers are Inferred and require human confirmation. Framework mappings are indicative. AEGIS does not certify ISO/IEC 42001 conformity, guarantee EU AI Act compliance, or replace legal advice.

Synthetic example — fictional data.